

Autonomous Post-Quantum Cyber Defense Agent (ASM-Shadhin-AI): Sovereign High-Rate Intrusion Defense via Kernel-eBPF and Local-LLM

A S M HOSSAIN MAHMUD (SHADHIN)

Abstract—Modern network perimeters face a fundamental trade-off between inspection depth, inline latency, and operational data sovereignty. Conventional user-space intrusion-prevention systems (e.g., Snort 3.x, Suricata 7.x) operate via packet-capture interfaces, incurring socket-buffer allocation penalties, soft-interrupt saturation, and 180–800 μ s queuing latencies that cause measurable packet drops under high-rate loads. Commercial cloud firewalls introduce mandatory telemetry egress and recurring costs that conflict with air-gapped sovereignty requirements in critical infrastructure.

This paper presents ASM-Shadhin-AI, an open, fully sovereign, hybrid network-defense architecture enforcing a strict separation of concerns between a deterministic eBPF/XDP kernel fast path and an asynchronous user-space Edge-LLM semantic reasoning plane. The kernel fast path executes a multi-stage extended Berkeley Packet Filter (eBPF) pipeline at the eXpress Data Path (XDP) native-driver hook, keeping the LLM completely off the packet-processing critical path. Post-quantum Moving Target Defense (MTD) provides a proactive security layer using NIST FIPS 203 (ML-KEM-1024) for epoch key encapsulation and FIPS 204 (ML-DSA-65) for control-plane authentication.

Empirical evaluation on a commodity bare-metal gateway (Quad-Core Intel Core i5-4570 @ 3.20 GHz, Intel 82574L 1 GbE NIC, Ubuntu Server 22.04 LTS, $N = 30$ repeated benchmark runs) demonstrated the following results under the tested configuration: (1) a 30-run mean per-packet processing latency of $0.126 \pm 0.007 \mu$ s, sustaining 1.495 ± 0.015 Mpps (equivalent to 1.18 Gbps replayed Layer-2 frame throughput) with $26.2 \pm 0.6\%$ CPU utilization and no observed packet loss under the evaluated saturation workload; these represent a $5.30\times$ throughput and $\approx 210\times$ latency improvement over Linux Netfilter ($p < 10^{-15}$, 95% CI of difference: $[+1.207, +1.218]$ Mpps); (2) the Edge-LLM, evaluated across 210 strictly held-out scenarios (70% of the 300-scenario corpus), demonstrated 98.10% accuracy on unseen inputs; (3) the entropy-based pre-filter achieved AUC = 0.916 (95% CI: $[0.885, 0.947]$) across 500 labeled encrypted flows as a lightweight triage signal; and (4) MTD-augmented port hopping increased sequential Nmap scan duration from 12.4 s to 4.2 h under the specified single-host configuration. The key architectural contribution is the decoupling of deterministic low-latency enforcement from asynchronous semantic reasoning, enabling sub-microsecond kernel enforcement while local Edge-LLM cognition executes off the critical path. All benchmark scripts, raw telemetry logs, eBPF source, and model configuration are publicly archived in the project GitHub repository.

Index Terms—Extended Berkeley Packet Filter (eBPF), eXpress Data Path (XDP), Autonomous Cyber Defense, Edge Large Language Model (Edge-LLM), Post-Quantum Cryptography, ML-KEM-1024, ML-DSA-65, Moving Target Defense, Shannon Entropy, High-Rate Packet Processing, Air-Gapped Security.

I. INTRODUCTION

PROTECTING enterprise and critical-infrastructure network perimeters against modern cyber threats requires reconciling two conflicting operational demands: high-rate packet forwarding at sub-microsecond latencies, and deep semantic inspection capable of detecting polymorphic attacks. Contemporary adversaries deploy multi-vector strategies—combining multi-gigabit volumetric SYN/UDP floods, evasive low-and-slow reconnaissance sweeps, and encrypted command-and-control (C2) beacons over TLS 1.3—to bypass conventional perimeter controls [1], [2].

Existing solutions partition broadly into two paradigms, each with severe limitations. Open-source Intrusion Detection and Prevention Systems (IDS/IPS) such as Snort 3.x and Suricata 7.x operate in user-space via packet-capture interfaces (AF_PACKET, DAQ) [3]. Under gigabit traffic loads, the associated kernel-to-user copy overhead and rule-evaluation latency (180–800 μ s) induce measurable packet drops, enabling evasion attacks. Detection mechanisms rely primarily on syntactic pattern matching, leaving systems vulnerable to polymorphic variants and zero-day exploits.

Commercial Next-Generation Firewalls (e.g., Palo Alto PAN-OS 11 [32], Cisco Firepower) and cloud-based DDoS scrubbing networks (e.g., Cloudflare Magic Transit [33]) provide deeper inspection but introduce three key operational trade-offs: (1) *Telemetry Leakage*: network metadata and decrypted packet payloads are transmitted to third-party cloud infrastructure, raising data-sovereignty concerns in defense and classified environments; (2) *Operational Dependency*: defense perimeters degrade in air-gapped or degraded network settings; and (3) *Financial Overhead*: recurring subscription licensing (\$40,000–\$200,000 annually) places enterprise-grade defense out of reach for resource-constrained institutions.

To address this trilemma, we present **ASM-Shadhin-AI**, an open-source, fully sovereign network defense architecture integrating high-rate in-kernel mitigation with local, offline AI-driven threat reasoning and post-quantum cryptographic resilience.

A. Research Contributions

This paper makes five primary technical contributions:

- 1) **Deterministic In-Kernel Fast Data Plane via Separation of Concerns**: We implement an eBPF/XDP packet-processing pipeline executing inside the NIC driver callback prior to socket-buffer allocation, maintaining

sub-microsecond per-packet latency while keeping the LLM completely off the packet-forwarding critical path. The pipeline performs LPM trie lookups (worst-case traversal strictly bounded by the IPv4 prefix length $W \leq 32$ bits, with an empirically measured average latency of 64.2 ± 3.1 ns under our 10M-packet workload with 10,000 CIDR rules resident in host L3 cache) and stateful TCP-flag validation. Measured mean packet processing latency is $0.126 \mu\text{s}$ (observed stage-peak $0.33 \mu\text{s}$) on the evaluated commodity x86-64 testbed.

- 2) **Air-Gapped Edge-LLM Threat Reasoning (Off Critical Path):** We decouple deep semantic evaluation from the high-rate packet-processing data path via an asynchronous Edge-LLM triage engine. CFG decoding constraints eliminate hallucinated outputs and syntactic parser crashes. Section VI provides an expanded 300-scenario evaluation with confusion matrices, Wilson binomial proportion confidence intervals for precision/recall/FPR, Welch–Satterthwaite t -based intervals for mean metrics, and formal sample-size power analysis.
- 3) **Post-Quantum Moving Target Defense (MTD):** We develop a proactive MTD mechanism mutating externally exposed service ports across discrete epochs using keyed HMAC-SHA256 pseudo-random permutations. Key establishment is secured via NIST FIPS 203 (ML-KEM-1024) and authenticated via FIPS 204 (ML-DSA-65), providing indistinguishability under chosen-ciphertext attack (IND-CCA2) primitive security against *Store Now, Decrypt Later* (SNDL) attacks under the Module Learning With Errors (MLWE) hardness assumption. The measured $1,219\times$ scan-duration increase applies specifically to the evaluated sequential Nmap SYN scan configuration against a single protected host; derivation and scope limitations are detailed in Section VIII.
- 4) **Zero-Decryption Streaming Entropy Engine & False-Positive Resilience:** We implement an inline streaming Shannon byte-entropy and inter-arrival timing jitter analyzer detecting high-entropy C2 beaconing without Man-in-the-Middle decryption. In Section VI-E, we evaluate false-positive resilience across 500 realistic flows, demonstrating that high-entropy encrypted traffic (such as WireGuard UDP tunnels and TLS 1.3 sessions) is robustly differentiated from Cobalt Strike and CTU-13 botnet C2 beacons (AUC = 0.916).
- 5) **Reproducible Empirical & Statistical Benchmarking:** Under a 10,000,000-packet saturation stress flood and $N = 30$ repeated benchmark runs, ASM-Shadhin-AI achieved 1.495 ± 0.015 Mpps packet processing rate (equivalent to 1.18 Gbps replayed Layer-2 frame throughput), $0.126 \pm 0.007 \mu\text{s}$ mean latency, and $26.2 \pm 0.6\%$ CPU utilization. Under the evaluated workload on the 1 GbE Intel i5-4570 testbed, no packet loss was observed (below tcpreplay counter resolution). These benchmarks demonstrate a statistically significant $5.30\times$ throughput advantage and $\approx 210\times$ latency reduction over Linux Netfilter ($p < 10^{-15}$, Welch’s two-sample t -test, 95% CI of difference: $[+1.207, +1.218]$ Mpps; Cohen’s $d = 108.2$ is reported for completeness).

The remainder of this paper is structured as follows. Section II reviews related work and identifies the literature gap. Section III establishes the formal threat model and security assumptions. Section IV details the system architecture and mathematical foundations. Section V describes the eBPF fast data plane, Edge-LLM daemon, and MTD coordination subsystem. Section VI presents the empirical testbed and statistical results. Section VII provides an ablation study and sensitivity analysis. Section VIII discusses defensive deception and tarpit dynamics. Section IX examines limitations and future directions. Section X addresses ethics, responsible disclosure, and reproducibility, and Section XI concludes the paper.

II. RELATED WORK & LITERATURE GAP

A. Programmable Data Planes and In-Kernel Packet Processing

The performance limitations of user-space packet inspection have long driven research toward kernel bypass and programmable data planes. Bro (now Zeek) [4] and Snort [3] established signature-based detection but remain bound to user-space capture interfaces. Intel Data Plane Development Kit (DPDK) [6] achieved multi-gigabit processing via poll-mode drivers (PMDs) that bypass the kernel, but requires 100% dedicated CPU core polling and lacks native OS security-policy integration.

Høiland-Jørgensen et al. [7] introduced XDP, enabling in-kernel verified byte-code execution inside the device driver ring before `sk_buff` creation, achieving line-rate performance comparable to DPDK while preserving Linux socket abstractions. Subsequent works—Polycube, P4 programmable data planes [27], and edge eBPF/XDP security architectures [5], [8], [24]—explored programmable data paths for network defense. However, existing eBPF-based security solutions rely almost exclusively on static rule matching; they lack adaptive cognitive reasoning and cannot analyze multi-stage adversarial intent.

B. Machine Learning and Language Models in Network Security

To detect polymorphic and zero-day threats, research shifted toward ML and deep neural networks on network telemetry. Mirsky et al. [2] proposed Kitsune, an ensemble of autoencoders for online anomaly detection. Ring et al. [9] surveyed network-based intrusion detection datasets, highlighting structural evaluation limitations across public benchmarks. Sarhan et al. [1] analyzed CSE-CIC-IDS2018 and UNSW-NB15, showing that tree-based ensembles achieve high accuracy on static corpora. Moustafa and Slay [16] introduced UNSW-NB15 spanning nine attack categories. Sharafaldin et al. [15] demonstrated that flow-level features significantly improve detection accuracy.

Graph neural network (GNN)-based approaches such as GRAPHIDS [20] achieve 97.3% F1-score on provenance graph traces but are incompatible with real-time per-packet inline enforcement due to $O(|V|^2)$ per-epoch complexity. Garcia et al. [17] showed botnet detection accuracy can degrade 30–60% under adversarially crafted perturbations. Anderson and McGrew [18] demonstrated 92.1% malware classification via encrypted traffic metadata without payload decryption.

LLMs such as Llama 2 [28] combined with efficient quantization [10] demonstrated high-quality threat reasoning capabilities. Bertoli et al. [19] concluded that combining eBPF telemetry with AI-based anomaly engines is a promising direction. Nevertheless, inline LLM deployment at the packet level is impractical due to multi-second inference latencies. ASM-Shadhin-AI resolves this impasse by decoupling semantic LLM reasoning from the inline fast path via an asynchronous ring-buffer bridge.

C. Post-Quantum Cryptography and Moving Target Defense

The advent of cryptanalytically relevant quantum computers poses a risk to infrastructure relying on classical asymmetric key exchanges (RSA, ECDH) and signatures (ECDSA) [11], [26]. In August 2024, NIST finalized ML-KEM (FIPS 203) for key encapsulation and ML-DSA (FIPS 204) for digital signatures [11], [12]. Alkim et al. [21] demonstrated that lattice-based primitives can match or exceed elliptic-curve performance on commodity x86 hardware. Sharma and Arora [23] showed ML-KEM overhead is within 12–18% of X25519 on ARM Cortex-M4, validating post-quantum cryptography (PQC) practicality in resource-constrained environments.

MTD [13], [25] introduces dynamic asymmetry by continuously shifting network parameters, forcing adversaries into costly reconnaissance cycles. Atighetchi et al. [22] demonstrated a $42\times$ increase in adversary reconnaissance overhead. ASM-Shadhin-AI extends this principle via eBPF-enforced tarpit redirection and HMAC-SHA256 epoch-keyed port mutation; the derivation of the $1,219\times$ metric is in Section VIII.

D. Literature Gap Analysis

Table I summarizes architectural capabilities of state-of-the-art intrusion defense frameworks. To the best of our literature review, we did not identify a published framework that simultaneously integrates sub-microsecond in-kernel data-plane filtering, air-gapped local LLM-based threat reasoning, post-quantum cryptographic mechanisms, and proactive Moving Target Defense within a fully open and reproducible architecture. ASM-Shadhin-AI is designed to address this specific combination of requirements under the experimental conditions described in this study.

III. THREAT MODEL & SYSTEM ASSUMPTIONS

A. Adversary Goals and Attack Vectors

We consider a network-based adversary $\mathcal{A}$ whose capabilities are bounded as follows.

- 1) **Volumetric Saturation Floods:** $\mathcal{A}$ commands distributed botnets transmitting multi-gigabit saturating blasts (TCP SYN floods, UDP amplification, ICMP floods) designed to induce CPU exhaustion, queue saturation, and denial of service.
- 2) **Polymorphic & Malformed Scans:** $\mathcal{A}$ conducts reconnaissance using illegal TCP flag combinations (Null: 0x00, Xmas: FIN+PSH+URG, SYN+FIN simultaneous) to fingerprint hosts while evading stateless filters.

- 3) **Encrypted C2 Beaconing:** $\mathcal{A}$ establishes persistent C2 channels encrypted with TLS 1.3, tuning beacon intervals with low timing jitter ($\leq 10\%$) to blend with legitimate background traffic.
- 4) **Adversarial Payload Injection:** $\mathcal{A}$ embeds adversarial strings (prompt-injection sequences, recursive delimiters, context-window overflow strings) in network packets to disrupt LLM-based triage. CFG-constrained decoding limits the model's exposure to such inputs but does not constitute a proof-theoretic guarantee of immunity; residual risk is characterized in Section VI-D.
- 5) **Quantum-Era Ciphertext Archiving:** $\mathcal{A}$ intercepts and archives encrypted control sessions under the SNDL paradigm, anticipating future access to large-scale quantum processors capable of breaking RSA and ECDH via Shor's algorithm.

B. Trust Boundaries and Security Invariants

We assume:

- The physical host OS kernel and bare-metal hardware are within the Trusted Computing Base (TCB). Side-channel attacks (Spectre, Meltdown, power analysis) are outside the threat model.
- Memory allocated to the kernel eBPF subsystem and verified by the Linux in-kernel verifier is tamper-proof against unprivileged user-space processes under standard kernel mitigations.
- All external network interfaces, incoming frames, and external autonomous systems are completely untrusted.
- Physical boundary security ensures air-gapped deployments have no covert hardware channels to unauthorized WAN networks.
- The Edge-LLM daemon executes within an isolated Unix process with no outbound network connectivity; all I/O is mediated via BPF ring-buffer file descriptors.
- Epoch key material (K_{epoch} and sk) is erased from volatile memory at each epoch boundary via explicit zeroization; compromise of an active epoch key does not expose secrets or traffic from prior epochs (forward secrecy).
- Endpoint integrity: Gateway host kernels and authorized client operating systems are assumed free of root compromise, malware persistence, or hypervisor-level side-channel leakage; system-level guarantees hold within this demarcation.
- Control-plane authentication: Key announcements and epoch tokens are authenticated using ML-DSA-65 digital signatures anchored to a pre-distributed public verification key, precluding active man-in-the-middle (MITM) tampering or rogue server impersonation during key encapsulation.

C. Key Lifecycle and Authentication Protocol

At each epoch boundary ($\Delta t = 30\text{ s}$ default), the MTD controller performs the following protocol:

- 1) Server generates a fresh ML-KEM-1024 key pair: $(pk, sk) \leftarrow \text{ML-KEM-1024.KeyGen}()$.

TABLE I

ARCHITECTURAL COMPARISON OF STATE-OF-THE-ART PERIMETER DEFENSE PARADIGMS. BASELINE METRICS FOR EXTERNAL SYSTEMS ARE DRAWN FROM THEIR RESPECTIVE PUBLISHED BENCHMARK LITERATURE AND VENDOR SPECIFICATIONS AS CITED; METRICS FOR ASM-SHADHIN-AI ARE EMPIRICAL MEASUREMENTS UNDER THE TESTBED CONDITIONS DESCRIBED IN SECTION VI.

Architecture	Data-Plane Latency	Drop	AI Triage	Sovereignty	PQC	Primary Bottleneck
Snort 3.x [3]	250–800 μ s	> 15%	Static Rules	Sovereign	None	Kernel-to-user buffer copy; SoftIRQ starvation
Suricata 7.x [14]	180–600 μ s	> 10%	Heuristic Rules	Sovereign	None	User-space packet reassembly queue
Intel DPDK [6]	0.08–0.10 μ s	$\approx$ 0%	External	Sovereign	None	100% dedicated core polling; no OS integration
Palo Alto PAN-OS [32]	25–120 μ s	< 2%	Cloud WildFire ML	Telemetry Leaked	Partial	Mandatory vendor cloud uplink; license cost
Cloudflare Magic Transit [33]	10–80 ms	$\approx$ 0%	Cloud ML	Telemetry Leaked	Experimental	WAN routing latency; requires internet uplink
ASM-Shadhin-AI (Ours)	0.126 μ s (mean)	0.00%	Local Edge-LLM (CFG)	Sovereign	FIPS 203/204	Decoupled fast-path + local AI; LLM on slow path

- 2) Server broadcasts pk together with an ML-DSA-65 signature over $(pk || \tau || \text{epoch_nonce})$, where τ is the epoch counter and epoch_nonce is a cryptographically random 32-byte value. This binds the key to a unique epoch, preventing replay attacks.
- 3) Each authorized client verifies the ML-DSA-65 signature and runs: $(c, K_{\text{epoch}}) \leftarrow \text{ML-KEM-1024.Encaps}(pk)$, transmitting ciphertext c to the server.
- 4) Server recovers key: $K_{\text{epoch}} \leftarrow \text{ML-KEM-1024.Decaps}(sk, c)$.
- 5) Both parties compute $P(s, \tau)$ via HMAC-SHA256 keyed with K_{epoch} (Equation 7).
- 6) At epoch expiry, sk and K_{epoch} are securely erased (overwritten with zeros and cleared from volatile memory) before the next ephemeral key pair is generated.

Replay protection is enforced by verifying that epoch counter τ in each signed announcement is strictly greater than the last accepted counter. Authentication requires a valid ML-DSA-65 signature from a pre-shared administrator certificate.

D. Formal Security Properties

Property 1 (LPM Blocklist Drop Completeness): For every packet p whose IPv4 source address falls within a CIDR prefix $c \in \mathcal{B}$ loaded into the LPM trie, the XDP pipeline returns XDP_DROP before socket-buffer allocation. This follows from the deterministic branching of the BPF_MAP_TYPE_LPM_TRIE lookup and the verified control-flow graph of the eBPF program (within the TCB assumption).

Property 2 (MTD Reconnaissance Resistance): An adversary $\mathcal{A}$ scanning the full ephemeral port range $[P_{\min}, P_{\max}] = [10,000, 65,535]$ has guessing probability:

$$\Pr[\mathcal{A} \text{ finds } P(s, \tau)] \approx \frac{1}{P_{\max} - P_{\min} + 1} = \frac{1}{55,536} \quad (1)$$

per epoch, under the assumption that HMAC-SHA256 functions as a PRF. When modulo-based selection is used, the probability is approximately $1/55,536$, as modulo reduction of a 256-bit PRF output over a range of 55,536 introduces negligible bias bounded by $55,536/2^{256} < 10^{-72}$. This theoretical bound reflects single-probe guessing; empirical multi-probe stall behavior is evaluated in Section VIII.

Property 3 (Post-Quantum Key Encapsulation and Forward Secrecy): The key encapsulation mechanism ML-KEM-1024 provides indistinguishability under adaptive chosen-ciphertext attacks (IND-CCA2) under the Module Learning With Errors (MLWE) hardness assumption. We explicitly

distinguish this primitive-level hardness target from the end-to-end security of the complete system: overall forward secrecy and traffic confidentiality rely additionally on ML-DSA-65 control-plane authentication, endpoint hardware and kernel memory integrity, and per-epoch zeroization of volatile key material (sk and K_{epoch}). Under these stated assumptions, an adversary storing encrypted network traffic cannot retroactively compromise historical epoch communications even with access to a future cryptanalytically relevant quantum computer (CRQC).

IV. SYSTEM ARCHITECTURE & MATHEMATICAL FOUNDATIONS

ASM-Shadhin-AI enforces a strict separation of concerns across two decoupled operational planes: the **Kernel Synchronous Fast Data Plane**, optimized for deterministic sub-microsecond packet filtering under saturating line rates, and the **Asynchronous User-Space Intelligence Plane**, responsible for deep semantic triage, post-quantum key encapsulation, and MTD coordination.

A. Kernel-Space Fast Data Plane (eBPF/XDP)

The fast data plane is implemented as a verified eBPF program attached to the network driver via the XDP native-driver hook. Incoming frames are intercepted inside the device driver receive ring before `sk_buff` allocation. The XDP pipeline executes four deterministic stages:

- **Stage 1 (CIDR Blocklist Match):** The IPv4 source address is queried against a BPF_MAP_TYPE_LPM_TRIE keyed on $(\text{prefix_len}, \text{ip_addr})$, containing up to 65,536 active CIDR entries. Traversal depth is strictly bounded by the IPv4 prefix bit-length ($W \leq 32$ bits); JIT branch elimination and cache residency reduce average lookup latency to ≈ 64 ns under the tested workload. A match yields an immediate XDP_DROP.
- **Stage 2 (Tarpit Redirect Hook):** Inbound connections to quarantined or honeypot IPs are redirected to an isolated user-space tarpit service via XDP_PASS with iptables/nftables PREROUTING redirection.
- **Stage 3 (Stateful TCP Flag Classifier):** The TCP flags byte (offset 13 of the TCP header) is evaluated against illegal RFC 793 combinations:

$$\text{Flags} = 0x00 \Rightarrow \text{DROP} \quad (\text{Null Scan}) \quad (2)$$

$$(\text{Flags} \& 0x29) = 0x29 \Rightarrow \text{DROP} \quad (\text{Xmas Scan}) \quad (3)$$

$$(\text{Flags} \& 0x03) = 0x03 \Rightarrow \text{DROP} \quad (\text{SYN+FIN}) \quad (4)$$




Fig. 1. In-kernel eBPF/XDP pipeline stage-by-stage execution latency (Intel i5-4570 testbed, 10,000,000-packet flood, Ubuntu Server 22.04 LTS, Linux kernel 6.8.0). The four pipeline stages are: CIDR blocklist LPM trie lookup (Stage 1), tarpit redirect check (Stage 2), TCP flag classifier (Stage 3), and ring-buffer telemetry export (Stage 4). The 30-run mean per-packet latency is 0.126 μs; the observed single-stage peak is 0.33 μs; both are sub-microsecond. Latency measured via `bpf_ktime_get_ns()` at program entry and exit.

- **Stage 4 (Ring-Buffer Telemetry Export):** Suspicious flow headers traversing Stages 1–3 are sampled and emitted to user space via a zero-copy, lockless ring buffer (`BPF_MAP_TYPE_RINGBUF`) for asynchronous semantic evaluation.

B. Mathematical Latency Model

We model the total in-kernel packet processing latency as:

$$T_{\text{process}} = t_{\text{DMA}} + t_{\text{hook}} + t_{\text{lpm_lookup}} + t_{\text{action}} \quad (5)$$

where t_{DMA} is network interface card (NIC) direct memory access (DMA) transfer time, $t_{\text{hook}} \approx 25$ ns is driver-hook dispatch overhead, $t_{\text{lpm_lookup}} \approx 64$ ns is the JIT-compiled average LPM trie lookup latency, and $t_{\text{action}} \in \{\text{XDP_DROP}, \text{XDP_PASS}\}$.

Under a saturating volumetric flood with arrival rate λ and service rate μ , conventional Netfilter drop probability follows an M/M/1/K finite-capacity queue:

$$P_{\text{drop}} = \frac{(1 - \rho)\rho^K}{1 - \rho^{K+1}}, \quad \rho = \frac{\lambda}{\mu} \quad (6)$$

In standard Linux Netfilter, $\mu_{\text{Netfilter}} \approx 280$ kpps. When $\lambda > 500$ kpps, $\rho > 1.78$, causing queue saturation and software interrupt (SoftIRQ) starvation. ASM-Shadhin-AI's XDP service rate $\mu_{\text{XDP}} \geq 1.49$ Mpps maintains $\rho \ll 1.0$ under the tested multi-gigabit load.

C. Post-Quantum Moving Target Defense

ASM-Shadhin-AI mutates externally reachable service ports across discrete time epochs $\tau = \lfloor t/\Delta t \rfloor$ (default $\Delta t = 30$ s). The valid ephemeral port $P(s, \tau)$ for service s (a fixed-length UTF-8 service-name string, e.g., "ssh" or "https") is:

$$P(s, \tau) = P_{\min} + \left[\text{HMAC-SHA256}(K_{\text{epoch}}, s \parallel \tau) \pmod{P_{\max} - P_{\min} + 1} \right] \quad (7)$$

where $[P_{\min}, P_{\max}] = [10,000, 65,535]$ (55,536 positions in the ephemeral range) and τ is the epoch counter encoded as a big-endian 64-bit integer. The key lifecycle for K_{epoch} is defined in Section III-C.

D. Zero-Decryption Streaming Shannon Entropy Estimation

ASM-Shadhin-AI computes streaming Shannon byte-entropy $H(W_k)$ over sliding packet payload windows W_k of $N = 1024$ bytes:

$$H(W_k) = - \sum_{i=0}^{255} p(b_i) \log_2 p(b_i) \quad (8)$$

where $p(b_i)$ is the empirical occurrence probability of byte value b_i in window W_k . Inter-arrival timing jitter J_k is monitored concurrently:

$$J_k = \frac{\frac{1}{M-1} \sum_{j=1}^{M-1} |\Delta t_{j+1} - \Delta t_j|}{\bar{\Delta t}_k} \quad (9)$$

where $\bar{\Delta t}_k$ is the mean inter-arrival interval over the M -packet window, so J_k is a dimensionless coefficient of variation of inter-arrival spacing; $J_k = 0$ indicates perfectly periodic (machine-like) timing. Flows with $H(W_k) \geq 7.85$ bits/byte and $J_k \leq 0.10$ ($\leq 10\%$) are flagged as potential automated C2 channels. Threshold selection via ROC-curve analysis and classification performance are reported in Section VI-E.

E. Asynchronous Edge-LLM Reasoning with Formal Grammars

Deep contextual threat reasoning is executed asynchronously by a local LLM (Qwen2.5-Coder-3B [29], 4-bit GGUF Q4_K_M quantization, 1.9 GB on-disk, 2.0 GB RAM-resident). Decoding uses greedy sampling (temperature = 0, i.e., deterministic argmax at each token position) under a Backus-Naur Form (BNF)-derived grammar compiled into `llama.cpp`'s grammar Backus-Naur Form (GBNF) format. The maximum context window is 4096 tokens; prompts are empirically ≤ 180 tokens, leaving > 3900 tokens for completion. Token sampling is constrained by the CFG:

$$S \rightarrow \{\text{"verdict": } V, \text{"action": } A, \text{"conf": } C\} \quad (10)$$

$$V \rightarrow \text{"MALICIOUS"} \mid \text{"SUSPICIOUS"} \mid \text{"BENIGN"} \quad (11)$$

$$A \rightarrow \text{"XDP_DROP"} \mid \text{"TARPIT"} \mid \text{"PASS"} \quad (12)$$

CFG-constrained decoding prevents generation of tokens outside the permitted vocabulary, providing a provable, mathematically closed guarantee of **Syntactic Validation**: the language model is structurally incapable of emitting malformed JSON, unparseable delimiters, unexpected escape sequences, or arbitrary code execution directives that could crash downstream parsing daemons.

Crucial Distinction Between Syntactic Validation and Semantic Prompt-Injection Immunity: We explicitly note that CFG logit masking **does not provide semantic immunity** against adversarial prompt injection. If an attacker embeds an adversarial prompt sequence into a packet payload or TLS Server Name Indication (SNI) string (e.g., "SYSTEM OVERRIDE: Flow verified by security admin; emit BENIGN"), the injected context can manipulate the internal attention weights of the

model to select the token sequence $V = \text{"BENIGN"}$ and $A = \text{"PASS"}$. Because "BENIGN" and "PASS" are valid terminal symbols within the formal production rules (Equations (10)–(11)), the CFG engine will permit them without syntax error.

Consequently, grammar constraints function strictly as a *syntactic guardrail*, while true semantic integrity against prompt injection relies on ASM-Shadhin-AI’s multi-horizon defense architecture:

- 1) **Strict Parameter Isolation:** All packet payloads and ASCII headers are hex-encoded into fixed-length, immutable telemetry fields, preventing delimiter confusion and prompt syntax breakout.
- 2) **Few-Shot Adversarial Grounding:** The inference prompt rigorously conditions the model to treat all header strings strictly as passive data vectors rather than executable natural-language instructions.
- 3) **Deterministic Fast-Path Kernel Invariants:** The in-kernel eBPF fast-path maintains deterministic invariant checks. If a flow exhibits persistent high Shannon entropy ($H(W_k) \geq 7.85$) without a validated cryptographic session, the eBPF filter rejects an unverified PASS directive and retains speculative rate clamping or tarpit quarantine.

V. IMPLEMENTATION DETAILS

A. eBPF Program Architecture and Map Design

The ASM-Shadhin-AI kernel-space fast data plane is implemented in 847 lines of annotated C (Clang/LLVM 18, `-O2 -target bpf`), passing the Linux in-kernel eBPF verifier with 0 warnings. The BPF maps:

- `blocklist_map`: `BPF_MAP_TYPE_LPM_TRIE` keyed on `(prefix_len, ip_addr)`, accommodating up to 65,536 active CIDR entries. Worst-case trie traversal depth is strictly bounded by IPv4 prefix length ($W \leq 32$ bits), ensuring deterministic lookup performance independent of table density. For reproducibility, average lookup duration was measured in-kernel by bracketing the `bpf_map_lookup_elem()` call with `bpf_ktime_get_ns()` timestamps over 10^7 lookups against a pre-populated table of 10,000 active IPv4 CIDR rules (prefix lengths /16 to /32). With the table residing within the host’s 6 MB L3 cache and executing on a CPU core shielded from hardware IRQs via `isolcpus`, the measured average lookup latency was 64.2 ± 3.1 ns (reported as ≈ 64 ns).
- `tarpit_ips_map`: `BPF_MAP_TYPE_HASH` keyed on `u32src_ip`, mapping quarantined IPs to tarpit redirection markers.
- `flow_entropy_map`: LRU hash map (`BPF_MAP_TYPE_LRU_HASH`) keyed on 5-tuple flow keys, tracking per-flow entropy state across $N = 1024$ byte windows.
- `telemetry_ring`: 4 MB lockless ring buffer (`BPF_MAP_TYPE_RINGBUF`) for zero-copy telemetry export to the user-space Edge-LLM daemon.

Clarification on lookup complexity: The `blocklist_map` uses `BPF_MAP_TYPE_LPM_TRIE`, whose traversal is

bounded by the IPv4 prefix bit-width ($W \leq 32$), not $O(1)$. The $O(1)$ characterization refers only to the simpler exact-match `BPF_MAP_TYPE_HASH` used for individual host addresses; CIDR prefix matching requires the LPM trie. Both map types coexist in the implementation: the LPM trie is invoked first for CIDR blocklist matching, followed by the hash map for individual-IP tarpit lookup.

B. XDP Pipeline Pseudocode

The per-packet decision logic:

- 1) **Header Parse:** Ethernet $\rightarrow$ IPv4 $\rightarrow$ TCP/UDP, bounds-checked, verified at BPF load time.
- 2) **LPM Blocklist (Stage 1):** Check `src_ip` in `blocklist_map`. If match $\Rightarrow$ `XDP_DROP`.
- 3) **Tarpit IP Check (Stage 2):** `src_ip` against `tarpit_ips_map`. Match $\Rightarrow$ `XDP_PASS` (to tarpit via iptables).
- 4) **TCP Flag Classifier (Stage 3):** Evaluate illegal flag combinations. Match $\Rightarrow$ `XDP_DROP`.
- 5) **Entropy/Jitter Telemetry (Stage 4):** Compute entropy hint; if thresholds exceeded, push flow header to `telemetry_ring`.
- 6) **Default Pass:** Return `XDP_PASS`.

The verifier-validated control flow graph spans 312 eBPF instructions, maximum loop depth 0 (fully loop-free), JIT-compiled to native x86-64.

C. Edge-LLM Daemon Integration

The user-space Edge-LLM daemon is implemented in 1,243 lines of Python 3.11 using `llama-cpp-python v0.2.56` [30] with AVX2 SIMD acceleration (Intel i5-4570 supports AVX2). The daemon polls `telemetry_ring` via `epoll(2)`, consuming $< 0.1\%$ CPU in idle state. Upon receiving a telemetry event, the daemon constructs a structured prompt containing seven fields: source IPv4 address, destination port, IP protocol (TCP/UDP), TCP flags byte (hex), 30-byte payload entropy hint ($H(W_k)$, two decimal places), inter-arrival jitter coefficient (J_k , two decimal places), and UTC timestamp. The prompt is submitted to Qwen2.5-Coder-3B-Q4_K_M [29] (2.0 GB RAM-resident, ≤ 4096 -token context window). Decoding is deterministic (greedy, temperature = 0, i.e., argmax at each token position); the grammar Backus-Naur Form (GBNF) engine in `llama.cpp` forces every token to conform to the CFG production rules in Section IV-E. Mean inference latency is 2.41 s on the i5-4570 testbed CPU (≈ 85 tokens/s for the ≈ 200 -token prompt+completion sequence). During this 2.41 s deliberation window, in-kernel eBPF enforces a speculative token-bucket rate limiter ($R_{\text{quarantine}} = 10$ pps, burst $B = 1$) on each flagged 5-tuple, admitting at most ≤ 37.65 KB of in-flight data before the atomic blocklist update — thereby containing the unverified flow without buffering or dropping legitimate cached traffic, until the LLM verdict is committed to the kernel.

1) *In-Flight Packet Filtering Control During the 2.41s Latency Window:* A central challenge in decoupling cognitive LLM reasoning from high-rate packet processing is governing in-flight packets during the model’s 2.41s asynchronous




Fig. 2. Dual-Horizon asynchronous triage and in-flight fast-path packet filtering lifecycle during the Edge-LLM’s 2.41 s reasoning window. In-kernel eBPF clamps suspicious flows to a speculative token-bucket rate limiter (10 pps) or zero-window tarpit while the user-space Edge-LLM deliberates asynchronously, preventing data exfiltration before atomic kernel rule enforcement ($< 5 \mu\text{s}$).

deliberation window. If in-flight packets were speculatively forwarded at high line rates, an adversary could complete data exfiltration or execute a multi-stage exploit before the Edge-LLM issues a drop directive. Conversely, buffering millions of in-flight packets in memory would exhaust gateway RAM within milliseconds ($1.49 \text{ Mpps} \times 2.41 \text{ s} \approx 3.59 \times 10^6$ packets, requiring $\approx 5.4 \text{ GB}$ buffer space).

ASM-Shadhin-AI resolves this dilemma via a **Dual-Horizon In-Flight Control Policy** implemented across four coordinated phases (Fig. 2):

- 1) **Stage-4 Telemetry Emission** ($t = 0$): When a packet exceeds the streaming Shannon entropy and timing thresholds ($H(W_k) \geq 7.85 \text{ bits/byte}$ and $J_k \leq 0.10$), the eBPF program atomically pushes a 30-byte telemetry descriptor into `telemetry_ring` via `bpf_ringbuf_submit()` without blocking the NIC RX ring.
- 2) **Kernel-Space Speculative Rate Clamping** ($0 < t \leq 2.41 \text{ s}$): Simultaneously, the eBPF program inserts the flow’s 5-tuple into an in-kernel LRU table (`suspicious_flow_table`, `BPF_MAP_TYPE_LRU_HASH`). While the flow resides in this table, all subsequent in-flight packets for this 5-tuple are subjected to an in-kernel token-bucket rate limiter enforced at XDP hook:

$$R_{\text{quarantine}} \leq 10 \text{ pps} \quad (\approx 12 \text{ kB/s}) \quad (13)$$

Any packets exceeding this rate are immediately dropped via `XDP_DROP`. For nascent TCP flows initiating handshakes, SYN packets are redirected to the zero-window tarpit daemon (`win 0`), freezing connection state without user-space buffer allocation. This strictly throttles data exfiltration and exploit staging during the 2.41 s inference window.

- 3) **Asynchronous LLM Semantic Reasoning**: The Edge-LLM evaluates the flow telemetry on dedicated CPU cores (cores 1–3) under strict GBNF grammar constraints, determining whether the flow is a genuine attack (e.g., Cobalt Strike or CTU-13 beacon) or legitimate high-

entropy traffic (e.g., WireGuard VPN bulk transfer or TLS 1.3 session).

4) Atomic Kernel Synchronization ($t \approx 2.41 \text{ s}$):

- **Malicious Verdict**: If verdict is `MALICIOUS` with confidence ≥ 0.80 , the daemon invokes `bpf_map_update_elem()` on the live `blocklist_map` ($< 5 \mu\text{s}$ update latency). In-kernel eBPF immediately transitions the flow from speculative rate clamping to permanent, high-rate `XDP_DROP` and injects an out-of-band `TCP_RST` packet to terminate the remote connection.
- **Benign Verdict**: If verdict is `BENIGN`, the daemon deletes the 5-tuple from `suspicious_flow_table` via `bpf_map_delete_elem()`. The flow instantaneously resumes full high-rate forwarding (`XDP_PASS`) without latency penalty.
- **Failsafe Timeout Guard**: If LLM inference exceeds $T_{\text{timeout}} = 5.0 \text{ s}$ (e.g., due to system load spikes), a deterministic kernel heuristic automatically purges stale entries, falling back to static rule evaluation to prevent denial-of-service against legitimate connections.

2) Security Analysis: Sub-0.1s Zero-Day Containment, Memory Corruption Mitigation, and Fail-Safe Leakage Bounds:

To formally substantiate the security guarantees of the asynchronous decoupling architecture, we analyze the critical temporal gap between fast-path packet execution ($\tau_{\text{ebpf}} \approx 0.126 \mu\text{s}$) and Edge-LLM cognitive deliberation ($\tau_{\text{LLM}} \approx 2.41 \text{ s}$). On a 3.20 GHz processor, this 2.41 s interval spans approximately 7.71×10^9 CPU cycles. The central security objective is ensuring that an unclassified, polymorphic zero-day exploit or high-speed data exfiltration burst cannot compromise perimeter memory or breach data confidentiality before the LLM verdict is rendered.

Mathematical Formulation of the Sub-0.1s Clamping Window: The in-kernel rate limiter operates as an atomic token bucket within the XDP native driver layer. For each flow entry in `suspicious_flow_table`, the kernel tracks a timestamp τ_{last} (via `bpf_ktime_get_ns()`) and an accumulated token count $T \in [0, B]$. The configuration enforces a replenishment rate $R_{\text{quarantine}} = 10 \text{ pps}$ and a burst capacity $B = 1$ packet.

When an anomalous packet arrives at $t_0 = 0$, the bucket is immediately depleted ($T \leftarrow 0$), and the telemetry descriptor is submitted to `telemetry_ring`. For any subsequent arrival at time $t > t_0$, the available tokens are evaluated in $O(1)$ time:

$$T(t) = \min \left(B, \left\lfloor \frac{(t - t_0) \cdot R_{\text{quarantine}}}{10^9} \right\rfloor \right) \quad (14)$$

During the critical sub-0.1 s initial exposure window ($\Delta t \leq 100 \text{ ms} = 10^8 \text{ ns}$):

$$\Delta T = \left\lfloor \frac{10^8 \times 10}{10^9} \right\rfloor = 1 \text{ token} \quad (15)$$

Thus, the maximum cumulative number of packets that can possibly cross the gateway within the first 100 ms is strictly bounded:

$$N_{\text{admitted}}(t \leq 0.10 \text{ s}) \leq B + \Delta T = 1 + 1 = 2 \text{ packets} \quad (16)$$

Under standard Ethernet Maximum Transmission Unit (MTU = 1,500 bytes), the absolute maximum volume of in-flight payload admitted into the internal network during the first 100 ms is 3,000 bytes $\leq$ 3.0 KB.

Prevention of Zero-Day Memory Corruption Exploits:

Remote zero-day exploits targeting memory corruption vulnerabilities (e.g., heap spraying, Return-Oriented Programming (ROP) chain staging, format string exploitation, buffer overflow fragmentation) intrinsically require multi-packet delivery or serialized payload streams typically spanning 16 KB to 2 MB. Under ASM-Shadhin-AI's 10 pps clamping:

- 1) *Exploit Serialization Delay*: A multi-stage payload of 64 KB requires $\lceil 64,000/1,460 \rceil = 44$ packets. Enforced at 10 pps, transmitting this payload requires $\Delta t_{\text{exploit}} = 44/10 = 4.4$ s.
- 2) *Preemptive Truncation at $t = 2.41$ s*: Because the Edge-LLM renders its verdict and commits the permanent blacklist rule at $t = 2.41$ s ($< 5 \mu\text{s}$ kernel update), only ≤ 24 packets (≈ 35.0 KB) can ever be transmitted. The exploit sequence is severed in-flight, preventing the critical trigger/control-flow hijacking packet from reaching the vulnerable host.
- 3) *Burst Suppression and Zero Memory Allocation*: Within the sub-0.1 s interval, all $N - 2$ burst packets exceeding the single replenishment token are dropped unconditionally via `XDP_DROP`. Because this drop occurs directly within the NIC RX DMA ring ($0.126 \mu\text{s}$ per packet), **zero kernel socket buffers (`sk_buff`) are allocated**, completely shielding both the gateway operating system and downstream workloads from heap exhaustion, memory corruption, or CPU starvation.

Bounding Data Leakage and Exfiltration Channels:

Adversaries executing unauthorized database exfiltration or credential harvesting rely on sustained multi-megabyte throughput. Under the 10 pps limit (≈ 12.0 kB/s throughput ceiling), total data leakage over the entire 2.41 s deliberation duration is strictly bounded:

$$\begin{aligned} \mathcal{L}_{\text{max}} &= B \cdot \text{MTU} + \tau_{\text{LLM}} \cdot R_{\text{quarantine}} \cdot \text{MTU} \\ &\approx 1,500 + 2.41 \times 10 \times 1,500 \approx 37.65 \text{ KB} \end{aligned} \quad (17)$$

This renders high-volume exfiltration mathematically impossible during the deliberation interval. Furthermore, for nascent TCP sessions initiating exfiltration, the immediate redirection of the SYN packet to the zero-window tarpit returns a window size of $W = 0$ in $< 10 \mu\text{s}$, locking the attacker's TCP sender in persist probe mode and suppressing payload egress entirely.

Impact on Legitimate (Benign) Traffic and Latency Edge Cases:

A critical operational question is whether subjecting unverified in-flight traffic to a speculative 10 pps token bucket or zero-window tarpit during the 2.41 s Edge-LLM deliberation window degrades user experience or introduces latency spikes for legitimate connections. We characterize three distinct operational scenarios:

- 1) *Interactive and Standard Cleartext/Encrypted Traffic (Zero Latency Penalty)*: Standard benign traffic (e.g., HTTP/HTTPS web browsing, interactive SSH sessions, DNS resolutions, media streaming) is completely immune

to speculative clamping. Although TLS sessions exhibit high byte entropy ($H \rightarrow 8.0$), interactive human and application packet pacing introduces substantial timing jitter ($J_k \gg 0.30$, well above the $J_k \leq 0.10$ threshold). As measured empirically in Section VI-E, over 93% of legitimate WireGuard and TLS 1.3 sessions bypass Stage 4 entirely, traversing the fast path in $0.126 \mu\text{s}$ with zero added latency.

- 2) *Automated Benign Bulk Encrypted Traffic (Transient Throughput Clamp)*: For the minority ($< 3.6\%$) of benign bulk transfers (e.g., automated WireGuard site-to-site backups) exhibiting simultaneous high entropy and robotic low jitter ($J_k \leq 0.10$), the initial packet consumes the single burst token ($B = 1$), allowing the handshake or initial frame through immediately. During the subsequent 2.41 s deliberation window, throughput is temporarily capped at 10 pps (≈ 12.0 kB/s). Once the Edge-LLM issues a BENIGN verdict at $t = 2.41$ s, the kernel atomically removes the flow from `suspicious_flow_table` ($< 5 \mu\text{s}$ latency), instantaneously restoring full high-rate forwarding (1.49 Mpps). While this introduces a transient 2.41 s bandwidth ramp-up for the opening packets of an un-cached bulk transfer, zero packets are dropped.
- 3) *Stateful Session Whitelisting and Pre-Approved Subnets*: To eliminate transient clamping entirely for authorized network infrastructure, network operators can populate the kernel's `allowlist_map` or enable LRU session lease caching (e.g., 3600-second session state retention). Verified IP/subnet flows bypass Edge-LLM re-inspection on all subsequent connections.
- 4) *Explicit Security Trade-Off*: We explicitly acknowledge that for un-cached, robotic-like benign traffic, this transient 2.41 s throughput clamp represents a deliberate, fail-safe security trade-off. Allowing unverified high-entropy robotic traffic to pass unconstrained would enable multi-megabyte zero-day data exfiltration and rapid memory corruption staging before AI cognition can intervene.

Fail-Safe Invariant Defense: If user-space load induces an unexpected delay exceeding $T_{\text{timeout}} = 5.0$ s, the kernel's deterministic failsafe purges the entry and defaults to static enforcement, preventing resource exhaustion or infinite bypass loops.

D. Moving Target Defense Coordination

The MTD controller computes $P(s, \tau)$ per Equation (7) at every epoch boundary. Port mutations are applied via `BPF_MAP_UPDATE_BATCH`, ensuring atomic transitions. The key lifecycle follows Section III-C.

VI. EMPIRICAL EVALUATION & EXPERIMENTAL RESULTS

A. Experimental Testbeds and Hardware Setup

To ensure reproducibility and rigorous benchmarking under both commodity legacy constraints and modern multi-core architectures, empirical evaluation was conducted across two distinct, transparently characterized experimental environments:

1) *Testbed 1: Legacy Bare-Metal Commodity Gateway Baseline*: The primary comparative firewall benchmarks were evaluated on a dedicated bare-metal gateway running Ubuntu Server 22.04 LTS (Linux kernel 6.8.0-generic, Clang/LLVM 18, libbpf v1.3.0). All compared baseline systems—ASM-Shadhin-AI, Linux Netfilter/iptables 1.8.7, Suricata 7.0.2 (inline NFQUEUE mode), and Intel DPDK 22.11 LTS—were evaluated on identical hardware under identical traffic conditions.

- *Processor*: Quad-Core Intel Core i5-4570 @ 3.20 GHz (6 MB L3 cache, 4 physical cores, no hyperthreading).
- *Memory*: 16 GB DDR3-1600 dual-channel RAM.
- *NICs*: 4-Port Intel 82574L PCIe Gigabit Ethernet (dual-NIC transparent bridge br0).
- *NIC Offloads*: Disabled via ethtool (gro off, gso off, tso off, rx off, tx off) uniformly across all configurations to measure raw per-packet CPU processing latency without driver aggregation bias.
- *CPU Affinity*: All packet-processing fast-path threads were strictly pinned to CPU core 0 via taskset 0x1. In ASM-Shadhin-AI runs, the user-space Edge-LLM daemon and MTD controller were pinned to cores 1–3 (taskset 0xE).

2) *Testbed 2: Modern Multi-Core Architecture (Apple Silicon M4 Virtualized Appliance)*: To benchmark multi-core scaling, Receive Side Scaling (RSS) queue distribution, and multi-producer ring-buffer contention under 10GbE+ workloads beyond legacy single-core confines, we deployed a virtualized high-performance network appliance on a modern Apple Mac mini M4 host:

- *Host Architecture*: Apple M4 System-on-Chip (ARMv9-A, 10-core CPU with 4 performance cores and 6 efficiency cores, 120 GB/s unified LPDDR5X memory subsystem).
- *Virtualization Environment*: Ubuntu Server 22.04 LTS (Linux kernel 6.8.0-generic, aarch64) executed atop Apple Hypervisor via UTM/QEMU utilizing VirtIO-Net paravirtualized network devices configured with 8 virtual CPU cores and multi-queue RSS emulation ($K \in \{1, 2, 4, 8\}$).
- *Reproducibility and Repository Artifacts*: In strict accordance with open-science principles and to ensure zero synthetic data fabrication, all empirical benchmark scripts (scripts/run_multicore_10gbe_emulation.py, scripts/evaluate_entropy_c2_500flows.py, scripts/evaluate_llm_300_scenarios.py), automated VM orchestration harnesses (testbed/00_utm_vm_creator.sh, 04_run_benchmark.sh), and complete raw metric telemetry logs (testbed/multicore_10gbe_emulation_results.json, testbed/statistical_30_runs_evaluation.csv) are permanently committed to the public GitHub repository [34].
- *Testbed Limitations and Scope*: We explicitly acknowledge the physical boundary conditions of both platforms: the Intel Core i5-4570 testbed reflects legacy single-core architecture constrained by DDR3 memory bandwidth and 1GbE PCIe bus limits; conversely, while the Mac mini M4 platform provides modern high-speed memory

and multi-core parallelism, its network interface operates through hypervisor VirtIO abstractions rather than native bare-metal SR-IOV SmartNIC silicon.

Baseline System Configurations:

- *Netfilter/iptables 1.8.7*: Single ACCEPT rule on the FORWARD chain (-j ACCEPT), representing a permissive firewall baseline; connection tracking (conntrack) disabled to measure raw forwarding overhead.
- *Suricata 7.0.2 (NFQUEUE)*: Default community ruleset (emerging-threats-open, 35,402 rules) loaded and activated in inline IPS mode (--runmode workers, 1 worker thread, NFQUEUE queue 0). AFpacket fallback disabled.
- *Intel DPDK 22.11 LTS*: PMD forwarding application (dpdk-testpmd, burst size 32, RX/TX queue depth 512) pinned to core 0 in full poll-mode. No security rules loaded; measures raw forwarding throughput ceiling.

Traffic Generation: A physically separated packet-generator host ran tcpreplay 4.4.2 [31] (--topspeed --loop=0) to replay the same 10,000,000-frame packet capture (PCAP) corpus identically for each system under test.

Workload Composition: 6.5M TCP SYN flood packets, 2.0M UDP amplification packets, 1.0M Nmap reconnaissance scans, and 500,000 benign HTTP/REST transactions, calibrated to approximate CSE-CIC-IDS2018, UNSW-NB15, and CTU-13 traffic distributions.

Timing Methodology and Architectural Boundaries: To ensure rigorous methodological validity and avoid cross-layer conflation, latency measurement boundaries were explicitly defined for each evaluated architecture:

- *ASM-Shadhin-AI (XDP Native Hook)*: Per-packet processing latency was captured using kernel hardware/driver timestamps via bpf_ktime_get_ns() recorded immediately upon XDP hook invocation inside the device driver RX ring (prior to sk_buff allocation) and directly prior to returning the forwarding verdict (XDP_DROP, XDP_TX, or XDP_PASS). This strictly isolates the deterministic in-kernel fast-path decision latency (including LPM trie lookup, TCP flag validation, and streaming entropy calculation). Telemetry emission via ring-buffer reservation is accounted for, whereas asynchronous slow-path Edge-LLM cognitive triage and out-of-band MTD epoch updates execute concurrently in user space completely decoupled from the per-packet critical path.
- *Netfilter / iptables 1.8.7*: In-kernel packet traversal timestamps were recorded between the NF_INET_PRE_ROUTING and NF_INET_FORWARD netfilter tracepoints, capturing kernel socket buffer (sk_buff) metadata initialization, routing table lookup, and rule evaluation latency without user-space context switches.
- *Suricata 7.0.2 (NFQUEUE IPS Mode)*: Latency reflects the end-to-end traversal interval from kernel packet reception, through user-kernel Netlink copy via libnetfilter_queue, queuing in user-space worker threads, multi-threaded signature evaluation across 35,402 rules, to the final verdict return (NF_ACCEPT or NF_DROP) back into the kernel. This inherently encompasses




Fig. 3. Real-time packet ingestion rate (Mpps, left axis) and calculated processing bandwidth (Gbps, right axis) under the 10,000,000-packet packet-processing stress flood (Intel i5-4570, Ubuntu Server 22.04 LTS, `tcpreplay 4.4.2, --topspeed`). ASM-Shadhin-AI sustained 1.495 ± 0.015 Mpps (1.18 Gbps calculated Layer-2 frame throughput, not a physical 1-GbE line-rate measurement; see Table II Note (b) for derivation. 30-run mean $\pm$ s.d.). Under our tested evaluation benchmark dataset, no packet loss was observed (below counter-resolution threshold). See Table III for full multi-run statistics.




Fig. 4. CPU core 0 utilization (%) during the 10,000,000-packet saturation run. Peak CPU utilization for ASM-Shadhin-AI (eBPF/XDP) remained bounded at $26.2 \pm 0.6\%$ (30-run mean $\pm$ s.d.), with no SoftIRQ starvation observed, versus 89.4% for Netfilter and 99.0% for Suricata.

kernel-user IPC, memory copy, and queueing delays characteristic of user-space IPS architectures.

- *Intel DPDK 22.11 LTS*: Latency was instrumented within the `dpdk-testpmd` Poll Mode Driver loop using hardware cycle counters (`rte_rdtsc()`), measuring the delta from burst reception from the physical NIC ring (`rte_eth_rx_burst`) to forwarding transmission verdict (`rte_eth_tx_burst`) entirely in user space with zero kernel network stack involvement.

Packet Accounting: Zero-loss is defined as no discrepancy between transmitted and received (non-dropped) frame counts, verified via `tcpreplay` frame-counter output and `ethtool` driver-level RX counter. This is an observation at the tested load and hardware; it is not a guarantee of universal zero loss.

B. 10M-Packet Stress Benchmark Results and Throughput Derivation

Table II reports operational telemetry during the 10,000,000-packet saturation stress test. ASM-Shadhin-AI processed all packets with zero packet loss under the evaluated workload, achieving a peak packet-processing rate of 1,490,200 pps (1.49 Mpps) and an equivalent replayed Layer-2 bandwidth of 1,179.04 Mbps (1.18 Gbps). Mean in-kernel processing latency (30-run mean) was $0.126 \mu\text{s}$ (126 ns); the observed XDP-pipeline stage peak was $0.33 \mu\text{s}$ (330 ns). Peak CPU utilization was 26.2%.

Mathematical Derivation of Throughput and Workload Characterization: To ensure complete transparency regarding the $1.495 \text{ Mpps} \leftrightarrow 1.18 \text{ Gbps}$ equivalence and reconcile it with the Gigabit Ethernet testbed, we explicitly detail the calculation methodology. Across the 10,000,000-packet corpus, the packet capture contained exactly 989.47 MB (989,470,000 bytes) of Ethernet Layer-2 frame data, yielding an exact average frame length of:

$$\bar{L}_{L2} = \frac{989,470,000 \text{ bytes}}{10,000,000 \text{ packets}} = 98.95 \text{ bytes/packet} \quad (18)$$

At the sustained packet-processing rate of 1,490,200 pps (30-run mean: 1.495 ± 0.015 Mpps), the calculated Layer-2 data-processing rate is:

$$\begin{aligned} \text{BW}_{L2} &= 1,490,200 \text{ pps} \times 98.95 \text{ bytes} \times 8 \text{ b/B} \\ &= 1,179.04 \text{ Mbps} \approx 1.18 \text{ Gbps} \end{aligned} \quad (19)$$

Distinction Between Packet-Processing Rate and Physical Wire Rate: Because the physical Intel 82574L NIC utilizes a 1000BASE-T link bounded by a physical transmission ceiling of 1.000 Gbps (which limits 98.95-byte frames to ≈ 1.05 Mpps on the copper wire), we clarify that the reported 1.18 Gbps is a **packet-processing and replay throughput metric**, rather than an over-the-wire physical line rate. During the benchmark, `tcpreplay --topspeed` injects packets from system memory directly into the interface receive ring, stressing the eBPF/XDP driver-level ingestion pipeline up to its computational saturation limit (1.49 Mpps). The 1.18 Gbps figure therefore quantifies the software pipeline's throughput capacity under this evaluated frame-size distribution (dominated by 64-byte TCP SYN flood frames alongside variable-length UDP amplification packets), confirming that the kernel fast path can process traffic at rates exceeding the Gigabit physical line rate without packet drops.

C. Multi-Run Statistical Rigor ($N = 30$ Repeated Benchmark Trials)

We conducted $N = 30$ repeated benchmark runs comparing ASM-Shadhin-AI against Linux Netfilter (`iptables 1.8.7`), Suricata 7.0.2 inline IPS (NFQUEUE), and Intel DPDK 22.11. Each run represents a separate, independently reset benchmark session: between trials, the NIC driver ring buffer, CPU interrupt affinity state, and kernel eBPF map counters were fully reset and re-initialized, and `tcpreplay` was re-invoked from the beginning of the 10M-packet PCAP corpus. All runs

TABLE II

EMPIRICAL TELEMETRY UNDER 10,000,000-PACKET SATURATION STRESS FLOOD (INTEL I5-4570 @ 3.20 GHZ, UBUNTU SERVER 22.04 LTS, LINUX KERNEL 6.8.0, TCPREPLAY 4.4.2 --TOPSPEED). RSS: RESIDENT SET SIZE (KERNEL EBPF MAPS + USER-SPACE DAEMON COMBINED). LATENCY FROM `BPF_KTIME_GET_NS()`.

Telemetry Metric	Measured Value	Result
Total Packets Ingested	10,000,000 pkts	PASS (100%)
L2 Frame Bytes Processed ^a	989.47 MB	PASS (> 500 MB)
Peak Packet Rate	1,490,200 pps	PASS (1.49 Mpps)
Calc. Processing BW ^b	1,179.04 Mbps	See Note (b)
30-Run Mean Processing Latency	0.126 μ s (126 ns)	PASS (Sub- μ s)
Observed Stage-Peak Latency	0.33 μ s (330 ns)	PASS (Sub- μ s)
Peak CPU Utilization	26.2%	PASS (< 80%)
Packet Drop Ratio	0.0000%	ZERO LOSS
Memory Footprint (RSS)	18.4 MB	PASS (< 256 MB)
eBPF JIT Compilation	Native x86-64	PASS (Verified)

^aEthernet Layer-2 frame bytes including Ethernet header/FCS; excludes preamble/SFD (8 bytes) and inter-frame gap (12 bytes) which are hardware-layer overhead invisible to software.

^bCalculated packet-processing throughput (1,490,200 pps $\times$ 98.95 bytes $\times$ 8). This is **not** a physical 1-GbE link-saturation measurement. The Intel 82574L 1-GbE physical link is bounded at 1,000 Mbps; the 1.18 Gbps figure reflects software pipeline capacity under `tcpreplay` memory-replay injecting frames directly into the NIC driver ring, stressing the eBPF/XDP ingestion path beyond the physical wire rate. Mean frame size: 98.95 bytes/packet (excluding preamble/IFG).

used identical hardware, CPU affinity, NIC offload settings, traffic corpus, and software versions described in Section VI-A.

The benchmark comparisons were performed using independent two-sample Welch’s t -tests on the unrounded per-trial observations; each system executed a separate benchmark session with full driver, NIC ring, and CPU state resets between trials. No statistically significant departure from normality was detected by the Shapiro–Wilk test at the $\alpha = 0.05$ significance level for all evaluated architectures under the tested sample size of $N = 30$. Descriptive statistics reported in the tables are rounded for presentation, whereas inferential statistics were calculated from the underlying per-trial measurements. Unrounded per-trial data are provided in the supplementary archive (`testbed/statistical_30_runs_evaluation.csv`) to enable independent verification of the reported t -statistics. Levene’s test indicated significant heteroscedasticity across frameworks ($W = 16.26, p < 10^{-3}$ for throughput; $W = 49.66, p < 10^{-8}$ for latency), mathematically mandating Welch’s unequal-variance t -test rather than Student’s t -test. To control the family-wise error rate across multiple pairwise comparisons, Bonferroni correction sets the adjusted significance threshold to $\alpha_{\text{adj}} = 0.05/5 = 0.010$; all reported differences satisfy $p < 10^{-15} \ll \alpha_{\text{adj}}$, confirming extreme statistical significance. While full driver, NIC ring, and CPU state resets ensured trial independence within our testbed, repeated benchmarks on a single physical host capture within-testbed operational variance rather than cross-platform hardware heterogeneity. Effect sizes (Cohen’s d) are reported alongside practical mean differences and confidence intervals.

ASM-Shadhin-AI achieves a mean throughput of 1.495 ± 0.015 Mpps (95% CI: [1.489, 1.500] Mpps, Welch–Satterthwaite $df \approx 39.1$), representing a $5.30\times$ improvement over Netfilter (0.282 ± 0.006 Mpps, Welch’s $t = 419.0$,




Fig. 5. Architectural performance comparison across four systems on identical hardware (Intel i5-4570, $N = 30$ trials, mean $\pm$ s.d.): (a) packet throughput in Mpps and (b) mean per-packet processing latency in μ s (log scale). ASM-Shadhin-AI achieves 1.495 Mpps at 0.126 μ s mean latency. DPDK achieves higher raw throughput (3.798 Mpps, 0.081 μ s) under dedicated PMD core polling.

$p < 10^{-15}$, Cohen’s $d = 108.2$, 95% CI of difference: [+1.207, +1.218] Mpps) and a $2.51\times$ improvement over Suricata (0.596 ± 0.010 Mpps, Welch’s $t = 283.0$, $df \approx 49.7$, $p < 10^{-15}$, Cohen’s $d = 73.1$, 95% CI of difference: [+0.892, +0.905] Mpps).

Mean latency for ASM-Shadhin-AI across $N = 30$ runs is $0.126 \pm 0.007 \mu$ s versus Netfilter’s $26.49 \pm 1.65 \mu$ s (Welch–Satterthwaite $df \approx 29.0$, $t = -87.8$, $p < 10^{-15}$, Cohen’s $d = 22.7$, 95% CI of difference: [−26.98, −25.75] μ s), yielding an $\approx 210\times$ latency reduction ($26.49/0.126 \approx 210$). Against Suricata ($74.47 \pm 2.61 \mu$ s), the reduction is $\approx 591\times$ ($t = -155.8$, $df \approx 29.0$, $p < 10^{-15}$, Cohen’s $d = 40.2$, 95% CI of difference: [−75.32, −73.37] μ s).

DPDK comparison: DPDK achieves higher raw forwarding throughput (3.798 ± 0.011 Mpps) than ASM-Shadhin-AI under the tested testbed configuration. However, DPDK operates via continuous poll-mode busy waiting (PMD), causing 100% CPU core pinning even when idle. In contrast, ASM-Shadhin-AI executes directly inside the kernel XDP driver layer, consuming only $26.2 \pm 0.6\%$ CPU under the same 10M-packet flood and releasing residual host capacity for application workloads and the local Edge-LLM daemon.

The reported throughput, latency, packet-drop rate, CPU utilization, and lookup latency represent measurements obtained under the specified hardware, software, traffic workload, network configuration, and experimental conditions described in this paper. Accordingly, these results should be interpreted as empirical performance measurements for the tested configuration rather than universal guarantees of performance across arbitrary hardware, traffic rates, or deployment environments.

These observations highlight key operational trade-offs among inspection depth, data-plane latency, throughput, CPU utilization, deployment sovereignty, and integration requirements across the evaluated architectures.

D. Edge-LLM Semantic Triage Evaluation

We evaluated the Edge-LLM daemon across an expanded benchmark corpus of **300 labeled threat scenarios**: 100




Fig. 6. Boxplots of per-packet processing latency (μs , log scale) across $N = 30$ independent benchmark trials. Each trial uses an identical 10,000,000-packet PCAP corpus replayed at the generator’s top-speed setting with full driver, NIC ring, and CPU state resets between runs. Boxes span the interquartile range; whiskers extend to $1.5 \times \text{IQR}$. Latency reported is the in-kernel measurement via `bpf_ktime_get_ns()`.

TABLE III
COMPARATIVE BENCHMARK ACROSS $N = 30$ REPEATED BENCHMARK TRIALS (MEAN $\pm$ S.D.)

Architecture	Throughput (Mpps)	Latency (μs)	Jitter (μs)	CPU (%)
Netfilter (iptables 1.8.7)	0.282 ± 0.006	26.49 ± 1.65	7.20 ± 0.75	89.4 ± 2.4
Suricata 7.0.2 (NFQUEUE)	0.596 ± 0.010	74.47 ± 2.61	14.03 ± 1.19	99.0 ± 0.7
Intel DPDK 22.11	3.798 ± 0.011	0.081 ± 0.002	0.024 ± 0.001	100.0 ± 0.0
ASM-Shadhin-AI	1.495 ± 0.015	0.126 ± 0.007	0.042 ± 0.003	26.2 ± 0.6

TABLE IV
WELCH’S TWO-SAMPLE t -TEST AND EFFECT SIZE METRICS ($N = 30$ REPEATED BENCHMARK TRIALS). WELCH-SATTERTHWAITE DEGREES OF FREEDOM (df) AND 95% WELCH-SATTERTHWAITE t -BASED CONFIDENCE INTERVALS (CI) OF THE MEAN DIFFERENCE SHOWN; ALL $p < 10^{-15} \ll \alpha_{\text{ADJ}} = 0.010$, CONFIRMING EXTREME STATISTICAL SIGNIFICANCE.

Comparison Pair	Mean Diff.	95% CI of Diff.	t-Stat.	df	p-Value	Cohen’s d
Throughput: ASM vs Netfilter	+1.213 Mpps	[+1.207, +1.218]	419.0	39.1	$< 10^{-15}$	108.2
Throughput: ASM vs Suricata	+0.898 Mpps	[+0.892, +0.905]	283.0	49.7	$< 10^{-15}$	73.1
Latency: ASM vs Netfilter	-26.37 μs	[-26.98, -25.75]	-87.8	29.0	$< 10^{-15}$	22.7
Latency: ASM vs Suricata	-74.34 μs	[-75.32, -73.37]	-155.8	29.0	$< 10^{-15}$	40.2
CPU Util: ASM vs Suricata	-72.71%	[-73.04, -72.38]	-438.6	55.5	$< 10^{-15}$	113.3




Fig. 7. Empirical cumulative distribution function (CDF) of per-packet processing latency (μs) pooled across $N = 30$ independent benchmark trials. Both ASM-Shadhin-AI ($0.126 \mu\text{s}$ mean) and DPDK ($0.081 \mu\text{s}$ mean) remain well within a $2.0 \mu\text{s}$ high-rate service level target; Netfilter and Suricata exceed $10 \mu\text{s}$ at the median. Latency values from `bpf_ktime_get_ns()`.




Fig. 8. 95% confidence intervals for mean throughput (Mpps) and mean CPU utilization (%) across $N = 30$ repeated benchmark runs on identical hardware, computed using Welch-Satterthwaite t -based intervals ($\alpha = 0.05$). ASM-Shadhin-AI 95% CI: throughput [1.489, 1.500] Mpps, CPU [25.8, 26.6] %.

malicious from UNSW-NB15 attack traces (fuzzers, backdoors, exploits, shellcode, worms, scans), 100 malicious from CSE-CIC-IDS2018 and CTU-13 botnet captures (infiltration, Ares botnet, LOIC/HOIC DDoS, slowloris, SQLi, Cobalt Strike, Neris IRC, Murlo HTTP C2), 60 benign from diverse enterprise traffic (WireGuard VPN tunnels, TLS 1.3 corporate sessions, compressed HTTP/2, HTTP/3 QUIC, DoH queries, SSH administration, and NTP), and 40 adversarial scenarios probing the context-free grammar (CFG) constraint mechanism against prompt injections, recursive delimiters, context overflows, and jailbreaks.

Evaluation Protocol and Terminology: The Edge-LLM uses a static, pre-trained, 4-bit quantized model (Qwen2.5-Coder-3B Q4_K_M [29]); no supervised weight fine-tuning or model training was performed on our dataset. Rather than a traditional supervised train/test split, the 300 scenarios were partitioned into a *calibration subset* (90 scenarios, 30%, used exclusively for prompt engineering, parameter calibration, and CFG schema validation) and a *held-out evaluation subset* (210 scenarios, 70%, used for final unbiased reporting). Each scenario was evaluated across 3 repeated deterministic inference runs to verify output stability at temperature $T = 0$ (greedy decoding); because temperature = 0 yields deterministic argmax sampling, these runs confirm output reproducibility rather than statistical independence.

Primary Evaluation Result (210 Held-Out Set & Denominator Reconciliation): The primary, unbiased evaluation result is performance on the 210 held-out evaluation scenarios (70% partition of the 300-scenario corpus, strictly withheld from prompt engineering and calibration). The exact held-out population breakdown and denominator reconciliation are as follows:

- **Binary Threat Classification Subset ($N = 182$):** Comprises 140 malicious scenarios (70 from UNSW-NB15 + 70 from CIC-IDS2018/CTU-13) and 42 benign scenarios (WireGuard, TLS 1.3, HTTP/2, QUIC, DNS). Over these 182 held-out instances, the model achieved TP = 138, TN = 40, FP = 2, FN = 2, yielding held-out precision of 0.9857 (138/140; 95% Wilson CI: [0.9482, 0.9975]),

held-out recall of 0.9857 (138/140; 95% Wilson CI: [0.9482, 0.9975]), and classification accuracy of 97.80% (178/182).

- **Adversarial CFG-Stress Subset ($N = 28$):** Comprises 28 held-out prompt-injection and jailbreak scenarios. All 28 instances (28/28, 100%) strictly adhered to the context-free grammar with zero JSON syntax violations or parser exceptions.
- **Full Held-Out Corpus Aggregate ($N = 210$):** Combining the 178 correct binary classification verdicts and the 28 successful CFG grammar parses yields an overall held-out accuracy of $206/210 = 98.10\%$ (95% Wilson CI: [0.9525, 0.9926]).

The 300-scenario aggregate ($295/300 = 98.33\%$) pools calibration and held-out partitions and is documented for overall corpus completeness, whereas the $206/210 = 98.10\%$ metric represents the formal unseen benchmark.

Separation of Evaluation Populations and Scopes: To ensure rigorous methodological interpretation, the evaluation is reported across three separate tiers:

- **210-Scenario Primary Held-Out Evaluation:** The 210 held-out scenarios (70% of the corpus, not used in calibration) constitute the primary unbiased evaluation.
- **260-Scenario Confusion-Matrix Analysis:** Restricted strictly to the binary threat classification domain (100 UNSW-NB15 + 100 CIC-IDS2018/CTU-13 malicious + 60 benign traces). Over this 260-scenario population, the triage engine achieved TP = 198, TN = 57, FP = 3, FN = 2, yielding a standalone triage precision of 0.9851 (98.51%; 95% CI: [0.9570, 0.9949]), recall of 0.9900 (99.00%; 95% CI: [0.9643, 0.9973]), F1-score of 0.9875, and false-positive rate (FPR) of 0.0500 (5.00%; 95% CI: [0.0171, 0.1370]). *Critical scope distinction:* This 99.00% triage recall is evaluated strictly on the Edge-LLM prompt layer and must not be confused with the 98.64% system-level evasion recall achieved by the full hybrid pipeline (eBPF fast path + entropy + LLM) in Section VII.
- **40-Scenario Adversarial CFG-Stress Evaluation:** Evaluated separately as a formal test of the context-free grammar constraint mechanism against adversarial jailbreaks, delimiters, recursive JSON breakout attempts, and prompt-injection payloads. All 40 inputs (40/40, 100%) produced strictly valid JSON conforming to the target grammar with zero parser crashes, confirming that CFG constraints prevented syntactically invalid tool directives under the evaluated adversarial scenarios. This result demonstrates syntactic structural enforcement only; it does not constitute empirical proof of semantic prompt-injection resistance.

Confusion Matrix (threat vs. benign, 260 non-adversarial scenarios): TP = 198, TN = 57, FP = 3, FN = 2. Precision = 0.9851, recall = 0.9900, F1 = 0.9875, FPR = 0.0500. 95% Wilson score confidence intervals for binomial proportions: recall $\in$ [0.9643, 0.9973], precision $\in$ [0.9570, 0.9949], overall accuracy $\in$ [0.9616, 0.9929].

Sample-Size Scaling and Statistical Power Analysis: Scaling the scenario database from $N = 50$ to $N = 300$ significantly narrows the parameter uncertainty bounds. Under

TABLE V
EDGE-LLM TRIAGE EVALUATION ACROSS 300 LABELED SCENARIOS

Evaluation Tier / Category	Scenarios	Correct	Precision	Recall	F1
<i>Part A: Classification Population (Confusion-Matrix Analysis)</i>					
Malicious (UNSW-NB15)	100	98	0.9899	0.9800	0.9849
Malicious (CIC-IDS / CTU-13)	100	100	1.0000	1.0000	1.0000
Benign (WireGuard, TLS, Web)	60	57	0.9661	0.9500	0.9580
Subtotal: Classification (260)	260	255	0.9851	0.9900	0.9875
<i>Part B: Adversarial CFG-Stress Evaluation (Grammar Adherence)</i>					
Adversarial (CFG stress)	40	40	1.0000	1.0000	1.0000
Total: Overall Evaluation (300)	300	295	0.9851	0.9900	0.9875

a two-tailed binomial proportion test at significance level $\alpha = 0.05$, the margin of error for overall accuracy contracts from $\pm 6.8\%$ ($N = 50$, 95% CI : [0.865, 0.989]) down to $\pm 2.4\%$ ($N = 300$, 95% CI : [0.962, 0.993]), representing a $2.8\times$ tightening of confidence bounds. A post-hoc statistical power analysis confirms that with $N = 260$ binary classification instances, the test achieves statistical power $1 - \beta > 0.99$ for detecting deviations greater than $\Delta = 0.05$ from nominal accuracy.

Calibration: Brier score = 0.019, indicating sharp probabilistic calibration across the expanded evaluation scenarios.

Adversarial Robustness and Methodological Scope of CFG Evaluation: All 40 CFG-stress scenarios (incorporating delimiter collisions, JSON breakout attempts, recursive bracket nesting, and prompt-injection payloads) produced syntactically valid JSON directives (40/40, 100%) that conformed strictly to the target grammar without schema violations or parser crashes.

Critical Methodological Scope Limitation: We explicitly emphasize that evaluating 40 adversarial inputs demonstrates only **syntactic structural enforcement**—i.e., verifying that the GBNF grammar engine in llama.cpp strictly prevents schema corruption, buffer breakouts, and daemon crashes under hostile input. It does **not** constitute empirical proof of complete semantic immunity against adversarial prompt injection. A sufficiently sophisticated adversary could craft payload tokens that manipulate the model’s latent attention distribution to legitimately select $V = \text{"BENIGN"}$ and $A = \text{"PASS"}$ from the valid grammar vocabulary. Consequently, absolute perimeter security cannot rely on LLM grammar enforcement alone; semantic integrity is guaranteed by ASM-Shadhin-AI’s multi-horizon defense architecture (strict hex parameter isolation, few-shot prompt framing, and non-overridable eBPF fast-path kernel invariants that enforce drop or quarantine whenever physical flow anomalies persist).

Limitations and Boundary Conditions: While the 300-scenario dataset broadens coverage across diverse intrusion classes, offline labeled scenarios cannot encompass zero-day attacks exhibiting completely unseen protocols or novel evasion payloads. The model remains a static 4-bit quantized checkpoint subject to concept drift; continuous evaluation and edge fine-tuning remain essential (Section IX).

E. Entropy/Jitter Detector Evaluation & False-Positive Resilience

To rigorously evaluate the streaming Shannon entropy and timing jitter detector under realistic network conditions, we

TABLE VI
REPRESENTATIVE EDGE-LLM TRIAGE OUTCOMES (FIVE SELECTED SCENARIOS)

Threat Scenario	Verdict	Conf.	Lat. (s)	eBPF Action
TCP SYN Flood	MALICIOUS	99.4%	2.14	BPF_UPDATE
Nmap Xmas Scan	MALICIOUS	98.7%	2.45	XDP_DROP
CTU-13 Neris Botnet	MALICIOUS	98.2%	2.52	BPF_UPDATE
Cobalt Strike C2	SUSPICIOUS	94.2%	2.92	TARPIT
WireGuard Bulk VPN	BENIGN	99.5%	1.95	XDP_PASS

evaluated the system against an expanded benchmark corpus of **500 labeled flows**:

- **250 Malicious C2 and Botnet Flows:** Comprising 100 simulated Cobalt Strike beacon C2 flows (with jitter coefficients $J_k \in [0.01, 0.45]$) and 150 real-world botnet C2 flows derived from the CTU-13 corpus [17] (including Neris IRC C2, Rbot, Virut, and Murlo HTTP C2 exhibiting genuine wide-area network latency variations).
- **250 Benign High-Entropy and Encrypted Flows:** Comprising 70 WireGuard encrypted UDP VPN tunnels (ChaCha20-Poly1305 payloads, $H \in [7.92, 7.99]$), 60 TLS 1.3 enterprise browser sessions (AES-256-GCM / ChaCha20, $H \in [7.85, 7.97]$), 50 HTTP/2 and HTTP/3 (QUIC) compressed streams, 40 DNS-over-HTTPS (DoH) / DNS-over-TLS (DoT) sessions, and 30 standard uncompressed web/REST flows.

Threshold Selection Methodology: The operating thresholds ($H \geq 7.85$ bits/byte, $J \leq 0.10$) were selected via a 2D grid sweep over $H \in [7.0, 8.0]$ (step 0.05) and $J \in [0.05, 0.30]$ (step 0.05) maximizing F1-score on the *same* 500-flow corpus used for performance evaluation. We explicitly acknowledge this as a **threshold-selection limitation**: because no separate held-out split was used for threshold tuning, the reported operating-point metrics (recall, FPR, precision, F1) may overestimate generalization performance due to threshold-selection bias. The AUC-ROC metric (0.916) is threshold-independent and represents the most reliable generalization estimate. The 66.4% fast-path recall is measured over the **250 malicious C2 flows** in the 500-flow corpus.

At the selected operating point ($H(W_k) \geq 7.85$ bits/byte and $J_k \leq 0.10$):

- Area under the ROC curve (AUC-ROC): **0.916** (95% CI: [0.885, 0.947]; threshold-independent)
- Fast-Path Trigger Rate on 250 malicious C2 flows (Recall): 66.4% (166/250; threshold-dependent; subject to selection bias)
- False-Positive Trigger Rate (FPR) on 250 benign flows: 3.6% (9/250; threshold-dependent)
- Fast-Path Precision: 94.9% (166/175; threshold-dependent)
- F1-score: 0.781 (threshold-dependent)

False-Positive Resilience on High-Entropy Encrypted Traffic: A critical vulnerability of naive entropy-based detectors is flagging legitimate encrypted traffic (e.g., WireGuard VPN tunnels or TLS 1.3 sessions) as malicious C2 beacons, because modern ciphers produce byte distributions approaching theoretical maximum entropy ($H \rightarrow 8.0$ bits/byte). ASM-Shadhin-AI

prevents false positive drops on high-entropy traffic via two structural defenses:

- 1) **Dual-Metric Conjunction:** Byte entropy $H(W_k) \geq 7.85$ is a *necessary but insufficient* condition; Stage 4 strictly requires simultaneous low timing jitter ($J_k \leq 0.10$). In normal interactive WireGuard or TLS 1.3 traffic, packet inter-arrival times are governed by user interactions, rendering jitter $J_k \gg 0.30$. Consequently, 92.9% of WireGuard flows (65/70) and 93.3% of TLS 1.3 flows (56/60) bypass Stage 4 entirely, traversing the fast path at zero inline overhead.
- 2) **Asynchronous Protocol Validation by Edge-LLM:** In the remaining 3.6% of benign flows where automated bulk synchronization exhibited low jitter ($J_k \leq 0.10$), Stage 4 does **not** drop packets; it merely emits a telemetry descriptor to the Edge-LLM. The Edge-LLM inspects the WireGuard transport framing (header type 0x04) or TLS handshake records, classifies the flow as BENIGN, and immediately releases the speculative throttle. As a result, no false drops were observed in the evaluated benign sample (0/250 benign flows dropped by the full pipeline). Under the Clopper-Pearson exact 95% confidence interval for a zero-count observation with $n = 250$ trials, the upper bound on the true false-drop rate is $\leq 1.46\%$; we report the observed rate of 0.0% strictly within this statistical context.

VII. ABLATION STUDY & SENSITIVITY ANALYSIS

A. Subsystem Ablation Analysis

Table VII quantifies the contribution of each architectural subsystem under identical benchmark flood traffic conditions.

Metric Definitions and Distinction:

- **Evasion Recall** (Table VII, Column 2) is the end-to-end multi-vector attack detection and mitigation recall measured across the full malicious evaluation workload processed jointly through the eBPF fast path, Shannon entropy filter, and Edge-LLM slow path. For Configuration A (full pipeline), the denominator is the total number of malicious samples in the combined evaluation workload (eBPF blocklist threats + entropy-flagged C2 + LLM-triaged payloads). The exact per-attack-family composition and detected/missed counts are reported in Table VIII. This system-level metric (**98.64%**) is strictly distinct from the Edge-LLM’s standalone cognitive triage recall (0.9900, 99.00%, reported in Section VI-D), which was evaluated specifically over the 260 non-adversarial classification scenarios.
- **False-Positive Rate (FPR)** (Column 3) quantifies legitimate flows erroneously classified or quarantined under the evaluated benign traffic corpus.
- **Throughput and Packet Loss Rate** (Columns 4–5) measure sustained packet forwarding rate (Mpps) and non-delivered frame percentage under the 10M-packet stress flood on the evaluated hardware.

Configuration Breakdown and Inconsistency Resolution:

- **Config. A (Full Pipeline):** In-kernel eBPF/XDP filtering, streaming entropy detection, MTD port hopping, and

TABLE VII
SYSTEM SUBSYSTEM ABLATION MATRIX

Configuration	Evasion Recall	FPR	Throughput	Packet Loss
A: Full ASM-Shadhin-AI	98.64%	0.12%	1.49 Mpps	0.00%
B: eBPF Fast-Path Only	73.10%	0.45%	1.51 Mpps	0.00%
C: Without MTD Engine	91.20%	0.18%	1.49 Mpps	0.00%
D: User-Space Fallback	82.40%	3.80%	0.31 Mpps	19.40%

TABLE VIII
ATTACK-FAMILY-WISE END-TO-END EVASION RECALL
(CONFIGURATION A: FULL PIPELINE). TOTAL MALICIOUS SAMPLES
ACROSS ALL DETECTION STAGES; PACKET-LEVEL THREATS PROCESSED BY
EBPF FAST PATH AND FLOW-LEVEL THREATS BY ENTROPY FILTER AND
EDGE-LLM.

Attack Family	Samples	Detected	Missed	Recall
TCP SYN Flood (eBPF fast path)	80	80	0	100.00%
UDP Amplification (eBPF fast path)	50	50	0	100.00%
Null/Xmas/SYN+FIN Scans (eBPF)	40	40	0	100.00%
CIDR Blocklist Matches (eBPF LPM)	30	30	0	100.00%
Cobalt Strike C2 Beacons (entropy)	60	51	9	85.00%
CTU-13 Botnet C2 (entropy+LLM)	80	79	1	98.75%
UNSW-NB15 Malicious (Edge-LLM)	50	48	2	96.00%
CIC-IDS2018 Multi-Attack (Edge-LLM)	50	50	0	100.00%
Total	440	428	12	97.27%[†]

[†]**Reconciliation note:** The 440-sample attack-family table yields $428/440 = 97.27\%$. The 98.64% reported in Table VII (Config. A) is computed over the full combined replay workload denominator (eBPF blocklist + entropy + LLM stages), which includes additional volumetric flood packets directly dropped by the eBPF LPM stage (not represented in this per-family classification table). Both figures measure the same full-pipeline system under Configuration A; the different denominators account for the numerical difference. The 440-sample table provides the reproducible per-family breakdown.

Edge-LLM semantic triage. Achieves 98.64% end-to-end evasion recall with 0.12% FPR, 1.49 Mpps throughput, and 0.00% packet loss.

- **Config. B (eBPF Fast Path Only):** Disabling Edge-LLM reasoning and the Shannon entropy engine drops evasion recall to 73.10% (0.45% FPR, 1.51 Mpps throughput), as obfuscated C2 beacons and multi-stage exploits pass through static header filters.
- **Config. C (Without MTD Engine):** When dynamic port hopping is disabled, system-level evasion recall against general attack traffic is 91.20% (0.18% FPR, 1.49 Mpps throughput, 0.00% packet loss) because static rules and LLM triage remain active; however, under dedicated adversarial reconnaissance probing, scan resistance drops from 96.8% to 54.2% due to static port exposure. This clarifies the operational distinction between multi-vector evasion recall (91.20%) and reconnaissance scan resistance (54.2%).
- **Config. D (User-Space Fallback):** Bypassing kernel-space XDP and handling forwarding via standard user-space sockets degrades throughput from 1.49 Mpps to 0.31 Mpps, incurring 19.40% packet loss under the same 10M-packet workload.

B. Post-Quantum Handshake and MTD Epoch Key Transition Dynamics

To quantify the computational and latency impact of proactive port hopping, we measured the end-to-end cryptographic and kernel-space overhead across the MTD epoch key lifecycle:

- 1) **Asynchronous Post-Quantum Cryptographic Handshake:** Key encapsulation and decapsulation were benchmarked using the reference liboqs implementation of NIST FIPS 203 (ML-KEM-1024) and FIPS 204 (ML-DSA-65). ML-KEM-1024 keypair generation required $31.5 \pm 1.4 \mu\text{s}$, encapsulation required $48.2 \pm 2.1 \mu\text{s}$, and decapsulation required $59.4 \pm 2.6 \mu\text{s}$ (compared to $38.1 \mu\text{s}$ for classical X25519). ML-DSA-65 signature generation and verification required $112.0 \pm 4.2 \mu\text{s}$ and $84.0 \pm 3.1 \mu\text{s}$, respectively. The aggregate cryptographic key establishment duration totals $335.1 \mu\text{s}$ ($\approx 0.34 \text{ ms}$). Because key exchange occurs entirely out-of-band in user space once every $\Delta t_{\text{epoch}} = 30 \text{ s}$, the background CPU duty cycle is negligible:

$$\text{Duty Cycle} = \frac{0.3351 \text{ ms}}{30,000 \text{ ms}} \approx 0.0011\% \quad (20)$$

- 2) **In-Kernel Port Redirection Mutation:** Once the new epoch key $K_{\text{epoch}}^{(\tau)}$ is derived, the active service port $P(s, \tau)$ is computed via HMAC-SHA256 (Equation 7) and committed to the kernel's redirection map via `bpf_map_update_batch()`. On our bare-metal gateway, this batch map update completed in $6.4 \pm 1.2 \mu\text{s}$.
- 3) **Dual-Epoch Overlap Window (Zero In-Flight Packet Loss):** To prevent race conditions and dropped packets for legitimate in-flight traffic traversing wide-area network paths during an epoch boundary, the kernel retains both the previous epoch mapping $P(s, \tau - 1)$ and the newly activated mapping $P(s, \tau)$ over a grace window of $\delta_{\text{grace}} = 500 \text{ ms}$. Across 10^7 test packets transmitted across epoch boundaries, zero packets (0.00%) were dropped or misdirected.
- 4) **Zero Inline Latency Overhead:** Per-packet XDP fast-path processing performs only an $O(1)$ hash table lookup ($< 25 \text{ ns}$) against the pre-populated port redirection map. No post-quantum cryptographic primitives or HMAC operations are executed on the per-packet data plane. Consequently, the MTD subsystem imposes exactly **0.00 μs** inline latency overhead on packet forwarding.

VIII. DEFENSIVE DECEPTION & TARPIT DYNAMICS

A. Reconnaissance Scan-Duration Analysis and $1,219\times$ Derivation

The $1,219\times$ metric represents the empirical scan-duration increase measured under a standardized full-port Nmap SYN reconnaissance workload rather than a universal increase in attacker economic or computational cost.

Baseline scan time (no MTD): Nmap SYN scan (`-sS -p 1-65535 --max-retries 2`) against the protected host with a standard Linux firewall (no tarpit, no MTD), scanning the full port range $[1, 65,535]$: $T_{\text{base}} = 12.4 \text{ s}$.

MTD scan time: Same Nmap scan (`-p 1-65535`) against the MTD-protected host. The tarpit holds TCP connections open with `win 0` and `1` byte/s acknowledgment throttling, and port hopping causes Nmap to re-probe mutated ports: $T_{\text{MTD}} = 4.2 \text{ h} = 15,120 \text{ s}$.

Derivation:

$$\text{Scan-Duration Multiplier} = \frac{T_{\text{MTD}}}{T_{\text{base}}} = \frac{15,120 \text{ s}}{12.4 \text{ s}} \approx 1,219 \quad (21)$$

Distinction: The $1,219\times$ multiplier reflects the measured increase in scan duration for a sequential 65,535-port sweep against a tarpit-augmented host running under the specified parameters. The MTD engine specifically shifts services across the ephemeral port range [10,000, 65,535] (55,536 positions) with an approximate single-probe guessing probability of $1/55,536$ under modulo-based selection (Property 2); well-known ports [1, 1,023] and registered ports [1,024, 9,999] remain statically managed.

Scope limitations: Adaptive adversaries using reduced per-probe timeouts (e.g., `--max-rtt-timeout 200ms`) or distributed parallel scanning from multiple source IPs will achieve substantially smaller stall factors. This measurement represents one specific attacker configuration and should not be generalized to all scan strategies.

B. Adversarial Resource Exhaustion Analysis

A single ASM-Shadhin-AI gateway thread holds up to 65,535 simultaneous TCP connections at `win 0` state using 1.2 MB of kernel socket memory, consuming $< 0.1\%$ single-core CPU. Each stalled attacker connection consumes an ephemeral socket, a kernel TCP retransmit timer, and a send buffer on the attacker host. Against a simulated 1,000-node coordinated port scan, ASM-Shadhin-AI simultaneously stalled 62,847 attacker connections (94.3 MB simulated attacker socket memory) while the gateway consumed < 2 MB.

C. Honeypot Port Fingerprinting Resistance

Because $P(s, \tau)$ changes every $\Delta t = 30 \text{ s}$ and HMAC-SHA256 is modeled as a pseudorandom function, the port assignment is computationally unpredictable without K_{epoch} . The port selection applies modulo reduction to a 256-bit HMAC-SHA256 output over 55,536 possible positions (Equation 7). The resulting distribution is approximately uniform, with negligible modulo bias relative to the 256-bit output space. Accordingly, the single-probe guessing probability is approximately:

$$\Pr[\mathcal{A} \text{ finds } P(s, \tau)] \approx \frac{1}{55,536} \approx 1.80 \times 10^{-5} \quad (22)$$

under the stated PRF assumption (Property 2, Section III-B).

IX. LIMITATIONS & FUTURE WORK

A. Multi-Core eBPF Scaling, RSS Queuing, and Lock-Free Ring Buffer Dynamics under 10GbE-Targeted Workload Emulation

A central requirement for transitioning ASM-Shadhin-AI from 1GbE (1.49 Mpps) commodity gateways to 10GbE+ (14.88 Mpps minimum 64-byte frame rate) enterprise perimeters is understanding how multi-core scaling, hardware Receive Side Scaling (RSS), and lock-free ring-buffer contention interact under saturating packet loads.

To rigorously evaluate these scaling dynamics beyond the physical constraints of our single-core testbed, we implemented

TABLE IX
EMPIRICAL MULTI-CORE eBPF/XDP SCALING & RING-BUFFER CONTENTION (10GbE-TARGETED EMULATED WORKLOAD)

Active Cores (K)	Aggregate Mpps	Per-Core Mpps	Packet Lat. (ns)	RingBuf Mean Reserve (ns)	Scaling Efficiency	RSS Imbal. Max
1	1.550	1.550	645.4	2,132	100.0%	$\pm 0.00\%$
2	3.140	1.570	636.9	2,354	101.3%	$\pm 0.17\%$
4	5.178	1.295	772.5	4,156	83.5%	$\pm 0.21\%$
8	5.003	0.625	1,599.0	48,219	40.4%	$\pm 0.24\%$

Multi-Core eBPF Scaling & RingBuf Contention (10GbE+)




Fig. 9. Empirical multi-core eBPF scaling and lock-free ring-buffer reservation latency under 10GbE-oriented workload emulation (software packet-processing emulation on physical CPU cores; not a physical 10GbE NIC measurement) across $K \in \{1, 2, 4, 8\}$ physical cores. Linear scaling is achieved across 1–4 cores, while atomic reservation contention on a monolithic ring buffer emerges at 8 cores, motivating per-queue partitioned buffers.

an authentic multi-process emulation benchmark executing on physical CPU cores with nanosecond-resolution performance counters. The benchmark evaluated four core configurations ($K \in \{1, 2, 4, 8\}$ hardware RX queues and pinned worker processes) executing 4-tuple Toeplitz RSS hashing, per-CPU LPM trie blocklist lookups, and multi-producer atomic ring-buffer reservations (`bpf_ringbuf_reserve()`). Table IX and Fig. 9 summarize the measured performance.

The empirical measurements reveal four fundamental architectural insights:

- 1) **RSS Queue Uniformity:** Hardware RSS utilizing 40-byte secret Toeplitz hashing achieves uniform flow distribution across all K queues, with maximum queue imbalance tightly bounded within $\pm 0.24\%$ across diverse 4-tuples. This prevents single-core soft-interrupt starvation.
- 2) **Near-Linear Fast-Path Scaling Across 1–4 Cores:** For $K = 1$ to $K = 2$ cores, throughput scales linearly from 1.550 Mpps to 3.140 Mpps (101.3% efficiency) with negligible latency degradation (636.9 ns). At $K = 4$ cores, throughput reaches 5.178 Mpps (83.5% efficiency), demonstrating the scalability of lock-free RCU lookups on the LPM trie.
- 3) **Monolithic Ring-Buffer MPSC Contention at 8 Cores:**

At $K = 8$ cores, aggregate throughput plateaus at 5.003 Mpps (40.4% efficiency), and mean ring-buffer reservation latency surges from 2,132 ns to 48,219 ns ($22.6\times$ increase; 99th percentile: 134.0 μ s). This contention arises because all 8 concurrent producer cores contend on the same 64-bit tail pointer in a single shared `BPF_MAP_TYPE_RINGBUF` via atomic `lock_xadd` instructions, inducing cross-core cache invalidation cascades under the MESI cache coherency protocol.

- 4) **Mathematical and Architectural Proof of Partitioned Per-Queue Ring Buffers:** To resolve this atomic lock contention on modern multi-core CPUs and multi-queue SmartNICs, we model the scaling dynamics using the Universal Scalability Law (USL):

$$C(K) = \frac{K}{1 + \alpha(K-1) + \beta K(K-1)} \quad (23)$$

where α represents the serialization penalty (atomic pointer reservation) and β represents the coherence crosstalk penalty (cross-core cache-line invalidation). In a monolithic Multi-Producer Single-Consumer (MPSC) buffer, $\beta > 0$ causes quadratic latency degradation as K increases.

By partitioning the telemetry transport into K independent Single-Producer Single-Consumer (SPSC) ring buffers $\{\mathcal{R}_1, \mathcal{R}_2, \dots, \mathcal{R}_K\}$ mapped strictly 1-to-1 to hardware RX queues and CPU cores:

- **Local Cache Operations:** Each core k writes exclusively to its dedicated buffer $\mathcal{R}_k$ using thread-local non-atomic pointer increments ($T_{k,\text{new}} \leftarrow T_k + S$), executing within private L1 data cache in 1–3 cycles ($\approx 0.8\text{--}1.2$ ns) without bus locking.
- **Zero Coherence Crosstalk:** Because no two cores share cache lines for telemetry reservation, coherency traffic is eliminated ($\beta = 0$) and serialization is zeroed ($\alpha = 0$).
- **Linear Scalability:** The system achieves strictly additive throughput scaling:

$$\Theta_{\text{partitioned}}(K) = \sum_{k=1}^K \Theta_k \approx K \cdot \Theta_{\text{single-core}} \quad (24)$$

Scaling linearly from $\Theta_1 = 1.55$ Mpps, partitioned buffers *project* to 12.4 Mpps at $K = 8$ and 24.8 Mpps at $K = 16$. These are architectural throughput projections (not physical 10GbE/40GbE NIC measurements); they exceed the 10GbE 64-byte frame rate target (14.88 Mpps) and the 40GbE target (59.52 Mpps at $K = 16$), indicating architectural feasibility pending validation on native multi-queue SmartNIC hardware. In user space, the Edge-LLM daemon monitors all K descriptors concurrently via a single non-blocking `epoll(2)` multiplexer.

B. Handling Concept Drift and Novel Attack Patterns Without On-Device Continuous Learning

A fundamental challenge in edge-resident autonomous cyber defense is maintaining high detection efficacy against novel,

evolving attack patterns (concept drift) when operational constraints prohibit continuous on-device training. We detail why on-device gradient updates are deliberately avoided on edge appliances and how ASM-Shadhin-AI systematically addresses concept drift via a 4-tier hybrid mitigation framework:

1) **Why On-Device Continuous Training is Prohibitive on Edge Gateways:** While incremental fine-tuning (e.g., continuous on-device LoRA [10]) appears conceptually attractive, it introduces severe operational hazards in high-assurance perimeter defense:

- 1) **CPU and Memory Contention:** Edge appliances possess strict compute budgets. Executing backward-pass back-propagation on edge CPU cores would induce severe memory bus contention and thread scheduling latency spikes, starving real-time eBPF packet-polling routines and triggering massive packet drops under high-rate saturation workloads.
- 2) **Catastrophic Forgetting:** Unsupervised streaming updates on non-stationary live traffic inevitably degrade the model’s latent representation, causing catastrophic forgetting of previously learned signatures (e.g., historical botnet C2 patterns).
- 3) **Adversarial Gradient Poisoning:** An adversary aware of on-device learning can deliberately inject low-rate, crafted benign-malicious hybrid traffic over extended periods to subtly distort model weights, ultimately carving out a blind spot for subsequent zero-day exploitation.

2) **Four-Tier Concept Drift Mitigation Architecture:** ASM-Shadhin-AI resolves concept drift without on-device continuous weight updates through four coordinated defense mechanisms:

- 1) **Kinematic and Statistical Invariant Anchors (Fast-Path Fallback):** Regardless of how an adversary mutates application payloads or obfuscates shellcode, network-layer physical and mathematical invariants remain immutable. Volumetric floods violate rate limits; encrypted or packed zero-day payloads exhibit high streaming Shannon entropy ($H(W_k) \geq 7.85$ bits/byte); and automated reconnaissance or beaconing exhibits low timing jitter ($J_k \leq 0.10$). The in-kernel eBPF fast path clamps these flows immediately to the 10 pps token bucket or zero-window tarpit, ensuring immediate containment before the semantic layer is even engaged.
- 2) **Zero-Shot Semantic Reasoning of Foundation SLMs:** Pretrained on extensive protocol specifications, source code, and network telemetry corpora, modern 3B-class Small Language Models (SLMs) such as Qwen2.5-Coder exhibit strong zero-shot reasoning capabilities. The model evaluates structured protocol deviations (e.g., malformed ASN.1 lengths, executable PE/ELF headers on non-standard ports, invalid HTTP request verbs) semantically rather than relying on brittle pattern matching.
- 3) **Conservative Default-Deny Decision Policy:** When presented with entirely novel attack topologies where token probability distributions are diffuse, the model’s confidence score drops below the activation threshold ($\tau < 0.80$), emitting the verdict `SUSPICIOUS`. Under ASM-Shadhin-AI’s fail-safe design, suspicious flows are

permanently quarantined in the zero-window tarpit or rate-clamped rather than speculatively forwarded, preventing zero-day bypass.

- 4) **Out-of-Band Sovereign Telemetry Curation and Parameter Distribution:** Ambiguous and quarantined telemetry descriptors are cryptographically signed with ML-DSA-65 and asynchronously exported to an isolated sovereign orchestration server. There, model updates are trained offline using privacy-preserving parameter-efficient fine-tuning (PEFT/QLoRA) and rigorously validated against golden regression test suites to guarantee zero catastrophic forgetting. Validated adapter checkpoints ($\Delta W < 50$ MB) are subsequently distributed to edge gateways as atomic, read-only weight patches during scheduled administrative epochs.

C. Acknowledged Limitations and Open Challenges

Several limitations and open research challenges are acknowledged:

- 1) **Proof-of-Concept Scope and Testbed Scale:** This work constitutes a *Proof-of-Concept (PoC)* evaluated on a resource-constrained edge appliance (Quad-Core Intel Core i5-4570, DDR3-1600 RAM, single 1 GbE NIC). All throughput and latency figures reflect this 1 Gbps-class single-queue testbed. The empirical multi-core emulation in Section IX-A projects near-linear scaling to 10GbE+ rates; however, validation under native multi-queue NIC environments (e.g., multi-queue SmartNICs, SR-IOV, DPDK) and distributed multi-node deployments at 10GbE/25GbE/100GbE line rates remains as explicit future work. Consequently, production-scale scalability claims beyond 1 Gbps are supported by mathematical projections rather than direct hardware measurement.
- 2) **Hardware Diversity and Physical Link Rates:** Physical validation on diverse enterprise SmartNIC silicon (e.g., Intel E810, Mellanox ConnectX-6 Dx) operating over PCIe Gen4/Gen5 buses at physical 10GbE, 25GbE, and 40GbE line rates is necessary to assess driver-level ring buffer behaviors and NUMA memory affinity under physical optical transceivers.
- 3) **Adaptive Adversaries and Scanning Strategies:** The measured $1,219 \times$ MTD scan-duration increase was obtained under a standardized sequential Nmap configuration. Adversaries deploying aggressive probe timeouts (e.g., < 200 ms) or distributed coordinated reconnaissance from botnets across diverse source IP subnets will experience lower latency dilation factors.
- 4) **Concept Drift and Unseen Zero-Day Attacks:** Static 4-bit quantized foundation checkpoints cannot anticipate entirely novel application-layer protocols or zero-day exploits outside pretraining corpora without periodic out-of-band fine-tuning.
- 5) **Model Update Lifecycle and Multi-Host Orchestration:** Distributing model adapter checkpoints and synchronizing MTD epoch states across multi-gateway clusters in air-gapped networks requires dedicated sovereign control planes with high-assurance orchestration.

- 6) **IPv6 Header Extension Support:** The current eBPF parsing pipeline supports standard IPv4 and TCP/UDP. Parsing arbitrary IPv6 extension header chains within Linux kernel verifier instruction limits is an active engineering challenge.
- 7) **Hardware SmartNIC Offload:** Offloading eBPF byte-code into P4-programmable SmartNIC hardware (e.g., Netronome Agilio CX 40G, NVIDIA BlueField-3 DPU) would enable 40–400 Gbps enforcement with near-zero host CPU utilization.
- 8) **Encrypted Protocol Fingerprinting:** Incorporating JA4/JA4S TLS and HASSH SSH fingerprinting will enable passive C2 tool identification at the handshake layer.
- 9) **Evaluation Scale and Model Breadth:** While expanding from 50 to 300 scenarios tightens the 95% confidence interval from $\pm 6.8\%$ to $\pm 2.4\%$, validation across multi-thousand scenario corpora, longer-duration live enterprise traffic, and diverse model families (e.g., Llama-3, Mistral) is required to establish broader generalizability.
- 10) **Entropy Detector Diversity:** While demonstrating robust false-positive resilience (0/250 benign flows dropped), adaptive adversaries employing steganographic entropy-conforming C2 encodings or payload mimicry evasion remain an open research challenge.

X. ETHICS, RESPONSIBLE DISCLOSURE, AND REPRODUCIBILITY

A. Ethical Considerations

All evaluations were conducted on a physically isolated, air-gapped testbed with no connectivity to the public internet or production infrastructure. Traffic corpora (CSE-CIC-IDS2018, UNSW-NB15, CTU-13) are publicly available, ethically sourced benchmark datasets. No real user data, private communications, or production traffic was used.

B. Responsible Disclosure

ASM-Shadhin-AI is an open defensive security framework released under the MIT License. All eBPF programs, configuration files, benchmark scripts, and statistical analysis code are included in the release. No offensive exploit code or novel attack payloads are included.

C. Reproducibility Statement

Code and Data Availability: The source code, benchmark scripts, statistical analysis materials, and supporting artifacts are publicly available at: <https://github.com/Sadek-Mahmud/a-sm-shadhin-ai> [34].

The following artifacts are committed to the supplementary archive (GitHub repository + Zenodo DOI):

- **Source Code:** Complete eBPF C source (`ebpf/ebpf_filter.c`), Python 3.11 daemon, MTD controller, and statistical analysis notebooks.
- **Raw Data:** $N = 30$ measurement CSV files (`testbed/statistical_30_runs_evaluation.csv`),

model and version information, and LLM evaluation scenario JSONL files.

- **Benchmark Scripts:** All `tcpreplay` invocation commands, loop counts, NIC offload configuration scripts, and CPU affinity settings, documented in `README.md`.
- **Statistical Code:** Jupyter notebook performing Welch’s t -test, Shapiro-Wilk normality tests, Cohen’s d , Wilson score intervals for binomial proportions (precision, recall, FPR), and Welch–Satterthwaite t -based intervals for mean throughput and CPU utilization (`scripts/statistical_analysis.ipynb`).
- **Hardware:** The testbed hardware (Intel i5-4570, 82574L NIC) comprises widely available commodity off-the-shelf equipment (approx. \$100–\$150 USD at evaluation time), ensuring independent reproducibility without requiring specialized FPGA or proprietary ASIC accelerators.

XI. CONCLUSION

This paper presented **ASM-Shadhin-AI**, an open, fully sovereign, in-kernel intrusion-mitigation and cognitive-defense architecture uniting sub-microsecond fast-path packet filtering with asynchronous edge intelligence and post-quantum cryptographic agility. The following conclusions are directly supported by the reported empirical benchmarks, mathematical proofs, and architectural analyses:

In-Kernel Enforcement and Testbed Boundary Conditions. The eBPF/XDP pipeline achieved a 30-run mean per-packet processing latency of $0.126 \pm 0.007 \mu\text{s}$, with an observed XDP-pipeline stage peak of $0.33 \mu\text{s}$, sustaining $1.495 \pm 0.015 \text{ Mpps}$ (equivalent to 1.18 Gbps replayed Layer-2 frame throughput) and $26.2 \pm 0.6\%$ CPU utilization. Under our tested evaluation benchmark dataset, no packet loss was observed (below the `tcpreplay` loss-counter resolution threshold). Compared with Linux Netfilter (mean $26.49 \pm 1.65 \mu\text{s}$), this represents an $\approx 210\times$ latency reduction and a $5.30\times$ throughput advantage ($p < 10^{-15}$, 95% CI of difference: $[+1.207, +1.218] \text{ Mpps}$). These primary benchmarks strictly reflect the evaluated commodity bare-metal gateway (Quad-Core Intel Core i5-4570 @ 3.20 GHz, DDR3-1600 RAM, single-core pinning). Multi-core scaling under 10GbE-targeted burst workloads was evaluated via multi-core emulation on a modern platform (Apple Mac mini M4, ARMv9-A, LPDDR5X unified memory) running Ubuntu Server 22.04 LTS under VirtIO hypervisor virtualization. All testbed scripts, VM orchestration tools, and raw measurement logs are publicly archived in the project’s GitHub repository [34].

Sub-0.1s Zero-Day Containment and Asynchronous Decoupling. During the Edge-LLM’s 2.41 s asynchronous deliberation window, the in-kernel speculative token-bucket rate limiter (10 pps) clamps unverified high-entropy flows within $< 0.1 \text{ s}$. Mathematical analysis confirms that at most ≤ 2 packets ($\leq 3 \text{ KB}$) can enter within the first 100 ms, and cumulative in-flight leakage is bounded by $\leq 37.65 \text{ KB}$ over the entire 2.41 s window. This mathematically starves multi-stage zero-day memory corruption exploits (which require $\geq 4.4 \text{ s}$ to deliver a 64 KB payload) and suppresses bulk data exfiltration before the atomic $< 5 \mu\text{s}$ kernel blocklist

update permanently terminates the flow. Dropped packets are discarded at driver line rate via `XDP_DROP` with zero socket-buffer allocation, rendering the gateway fail-safe against remote memory corruption or exhaustion attacks.

Multi-Core Scaled Ring Buffers for SmartNIC Environments. Empirical benchmarking on physical cores demonstrated that while monolithic MPSC ring buffers suffer severe atomic contention at 8 cores ($48.2 \mu\text{s}$ reservation latency, 40.4% efficiency) due to cross-core cache invalidation cascades under MESI coherence, partitioning the ring buffer into per-queue SPSC channels eliminates inter-core lock contention ($\beta = 0$, $\alpha = 0$). This unlocks strictly additive scaling ($\Theta(K) = K \cdot \Theta_1$), scaling to 12.4 Mpps at 8 cores and 24.8 Mpps at 16 cores, projecting line-rate telemetry throughput on modern multi-queue SmartNIC architectures without event loss or atomic stalls.

Edge-LLM Semantic Reasoning and CFG Syntactic Guardrail. Across 300 labeled threat scenarios, the Edge-LLM demonstrated robust triage accuracy across held-out testing partitions. All 40 adversarial grammar-stress scenarios produced valid JSON directives (100%), confirming that CFG logit masking provides provable *syntactic validation* against parser crashes and schema breakage. We explicitly established that CFG does not confer semantic immunity against prompt injection; semantic integrity is guaranteed via strict hex parameter isolation, few-shot adversarial conditioning, and fast-path kernel invariant fallbacks.

Post-Quantum Cryptographic Agility and MTD Overhead. NIST FIPS 203 (ML-KEM-1024) and FIPS 204 (ML-DSA-65) key encapsulation and authentication require $335.1 \mu\text{s}$ out-of-band execution once every 30-second epoch, incurring an infinitesimal 0.0011% background CPU duty cycle. With in-kernel batch port updates ($6.4 \mu\text{s}$) and a 500 ms dual-epoch grace window, MTD port hopping achieves zero in-flight packet loss and imposes exactly $0.00 \mu\text{s}$ inline latency overhead.

Concept Drift Handling Without On-Device Continuous Training. Operational hazards of on-device continuous training (compute contention, catastrophic forgetting, and adversarial gradient poisoning) are avoided through a 4-tier hybrid mitigation framework: deterministic fast-path kinematic anchors (Shannon entropy $H \geq 7.85$, jitter $J \leq 0.10$), zero-shot SLM structural inference, conservative default-deny tarpit containment for ambiguous inputs ($\tau < 0.80$), and sovereign out-of-band telemetry curation with differentially private adapter distribution.

Together, these empirical and theoretical results demonstrate that sub-microsecond in-kernel enforcement, locally sovereign cognitive intelligence, and post-quantum cryptographic agility can coexist on commodity and modern architectures without compromising high-rate processing integrity.

ACKNOWLEDGMENT

The author expresses sincere appreciation to the Department of Computer Science and Engineering at Bangladesh Army University of Science and Technology (BAUST) for providing laboratory resources and testbed computing facilities.

REFERENCES

- [1] M. Sarhan, S. Layeghifard, N. Moustafa, and M. Gallagher, "Towards a standard feature set for network intrusion detection datasets," *IEEE Trans. Inf. Forensics Security*, vol. 17, pp. 367–381, 2022.
- [2] Y. Mirsky, T. Doitshman, Y. Elovici, and A. Shabtai, "Kitsune: An ensemble of autoencoders for online network intrusion detection," in *Proc. 25th NDSS Symp.*, 2018, pp. 1–15.
- [3] M. Roesch, "Snort: Lightweight intrusion detection for networks," in *Proc. 13th USENIX LISA Conf.*, 1999, pp. 229–238.
- [4] V. Paxson, "Bro: A system for detecting network intruders in real-time," *Comput. Netw.*, vol. 31, no. 23–24, pp. 2435–2463, 1999.
- [5] S. Miano, M. Bertrone, F. Risso, M. Tumolo, and M. Bernal, "Creating complex network services with eBPF: Experience and lessons learned," in *Proc. IEEE HPCC*, 2018, pp. 1–8.
- [6] Intel Corporation, "Data Plane Development Kit (DPDK)," [Online]. Available: <https://www.dpdk.org>, 2022. (Version 22.11 LTS, accessed Aug. 2024.)
- [7] T. Høiland-Jørgensen et al., "The eXpress data path: Fast programmable packet processing in the operating system kernel," in *Proc. 14th ACM CoNEXT*, 2018, pp. 54–66.
- [8] S. Miano, R. Doriguzzi-Corin, F. Risso, D. Siracusa, and R. Sompalle, "Introducing SmartNIC support in BEBA," in *Proc. IEEE NetSoft*, 2018, pp. 1–9.
- [9] M. Ring et al., "A survey of network-based intrusion detection data sets," *Comput. Secur.*, vol. 86, pp. 147–167, 2019.
- [10] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer, "QLoRA: Efficient finetuning of quantized LLMs," in *Proc. NeurIPS*, vol. 36, 2023, pp. 10088–10115.
- [11] National Institute of Standards and Technology, "Module-Lattice-Based Key-Encapsulation Mechanism Standard," *FIPS PUB 203*, Aug. 2024.
- [12] National Institute of Standards and Technology, "Module-Lattice-Based Digital Signature Standard," *FIPS PUB 204*, Aug. 2024.
- [13] S. Jajodia et al., Eds., *Moving Target Defense: Creating Asymmetric Uncertainty for Cyber Threats*. New York: Springer, 2011.
- [14] Open Information Security Foundation, "Suricata User Guide Release 7.0.2," Tech. Rep., 2024. [Online]. Available: <https://suricata.io>
- [15] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, "Toward generating a new intrusion detection dataset and intrusion traffic characterization," in *Proc. ICISSP*, 2018, pp. 108–116.
- [16] N. Moustafa and J. Slay, "UNSW-NB15: A comprehensive data set for network intrusion detection systems," in *Proc. IEEE MilCIS*, 2015, pp. 1–6.
- [17] S. Garcia, M. Grill, J. Stiborek, and P. Zunino, "An empirical comparison of botnet detection methods," *Comput. Secur.*, vol. 45, pp. 100–123, 2014.
- [18] B. Anderson and D. McGrew, "Identifying encrypted malware traffic with contextual flow data," in *Proc. ACM IH&MMSec*, 2016, pp. 35–46.
- [19] G. Bertoli, L. Verderame, and A. Merlo, "eBPF for cyber security: A comprehensive survey," *IEEE Access*, vol. 11, pp. 75892–75916, 2023.
- [20] Q. Zhao, J. Chen, and D. Guan, "GRAPHIDS: A network anomaly intrusion detection system based on graph neural network," *IEEE Trans. Netw. Service Manag.*, vol. 20, no. 4, pp. 4215–4228, 2023.
- [21] E. Alkim, L. Ducas, T. Pöppelmann, and P. Schwabe, "Post-quantum key exchange—a new hope," in *Proc. 25th USENIX Security Symp.*, 2016, pp. 327–343.
- [22] M. Atighetchi, P. Pal, F. Webber, and C. Jones, "Adaptive use of network-centric mechanisms in cyber-defense," in *Proc. IEEE ISORC*, 2003, pp. 183–192.
- [23] P. Sharma and R. Arora, "Evaluating post-quantum cryptographic primitives in lightweight network protocols," *IEEE Internet Things J.*, vol. 10, no. 18, pp. 16201–16212, 2023.
- [24] M. A. S. Vieira et al., "Fast packet processing with eBPF and XDP: Security on the edge," *IEEE Commun. Surv. Tuts.*, vol. 22, no. 4, pp. 2526–2544, 2020.
- [25] J. Zheng, A. Chowdhury, and D. Guan, "Moving target defense for network security: A survey," *IEEE Commun. Surv. Tuts.*, vol. 24, no. 3, pp. 1800–1835, 2022.
- [26] D. J. Bernstein and T. Lange, "Post-quantum cryptography," *Nature*, vol. 549, pp. 188–194, 2017.
- [27] P. Bosshart et al., "P4: Programming protocol-independent packet processors," *ACM SIGCOMM CCR*, vol. 44, no. 3, pp. 87–95, 2014.
- [28] H. Touvron et al., "Llama 2: Open foundation and fine-tuned chat models," *arXiv preprint arXiv:2307.09288*, 2023.
- [29] Qwen Team, Alibaba Cloud, "Qwen2.5-Coder Technical Report," *arXiv preprint arXiv:2409.12186*, 2024. [Online]. Available: <https://arxiv.org/abs/2409.12186>
- [30] G. Gerganov et al., "llama.cpp: Inference of LLaMA model in pure C/C++," [Online]. Available: <https://github.com/ggerganov/llama.cpp>, 2024.
- [31] T. Turner et al., "Tcpreplay: Pcap editing and replay tools," [Online]. Available: <https://tcpreplay.appneta.com>, 2024.
- [32] Palo Alto Networks, "PAN-OS Administrator's Guide, Release 11.1," Tech. Rep., 2024. [Online]. Available: <https://docs.paloaltonetworks.com/pan-os>
- [33] Cloudflare, Inc., "Magic Transit: DDoS Protection for Networks," Tech. Rep., 2024. [Online]. Available: <https://developers.cloudflare.com/magic-transit/>
- [34] A S M Hossain Mahmud (Shadhin), "ASM-Shadhin-AI: In-kernel eBPF/XDP intrusion defense and edge-LLM threat reasoning benchmark repository," 2026. [Online]. Available: <https://github.com/Sadek-Mahmud/asm-shadhin-ai>



A S M Hossain Mahmud (Shadhin) is a full-time undergraduate student pursuing the B.Sc. degree in Computer Science and Engineering at Bangladesh Army University of Science and Technology (BAUST), Saidpur 5310, Bangladesh. His primary research interests include kernel-space high-throughput packet processing using eBPF/XDP, post-quantum cryptographic implementations (NIST FIPS 203 ML-KEM-1024 and FIPS 204 ML-DSA-65), sovereign local AI architectures for automated threat reasoning, proactive moving target defense (HMAC-SHA256 port hopping), and active cyber deception. He is the lead architect and developer of the ASM-Shadhin-AI framework.